

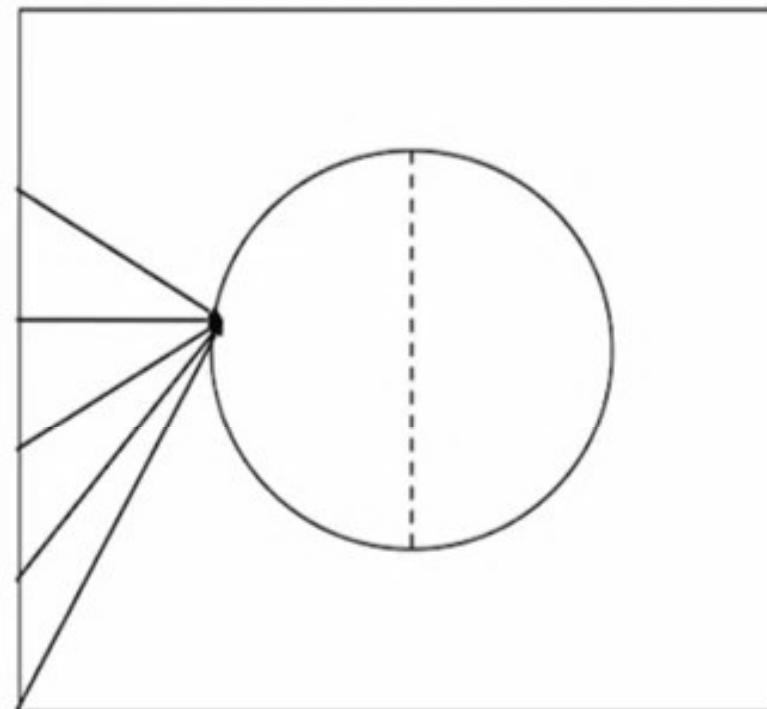


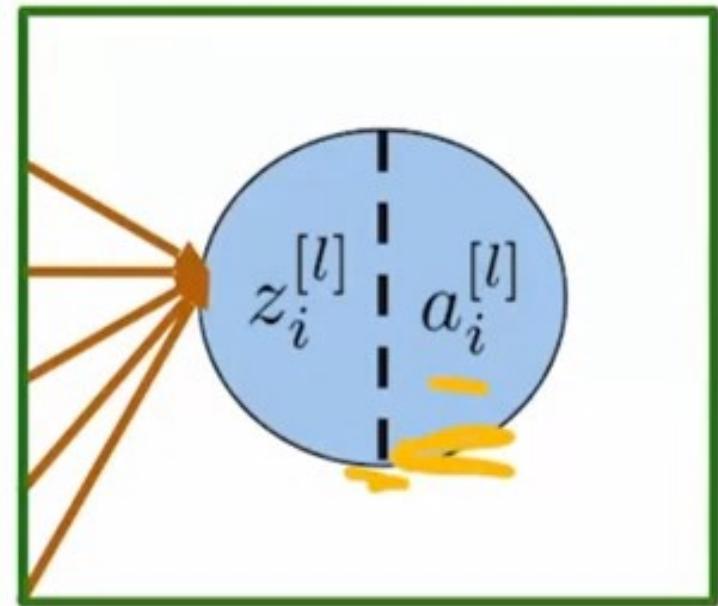
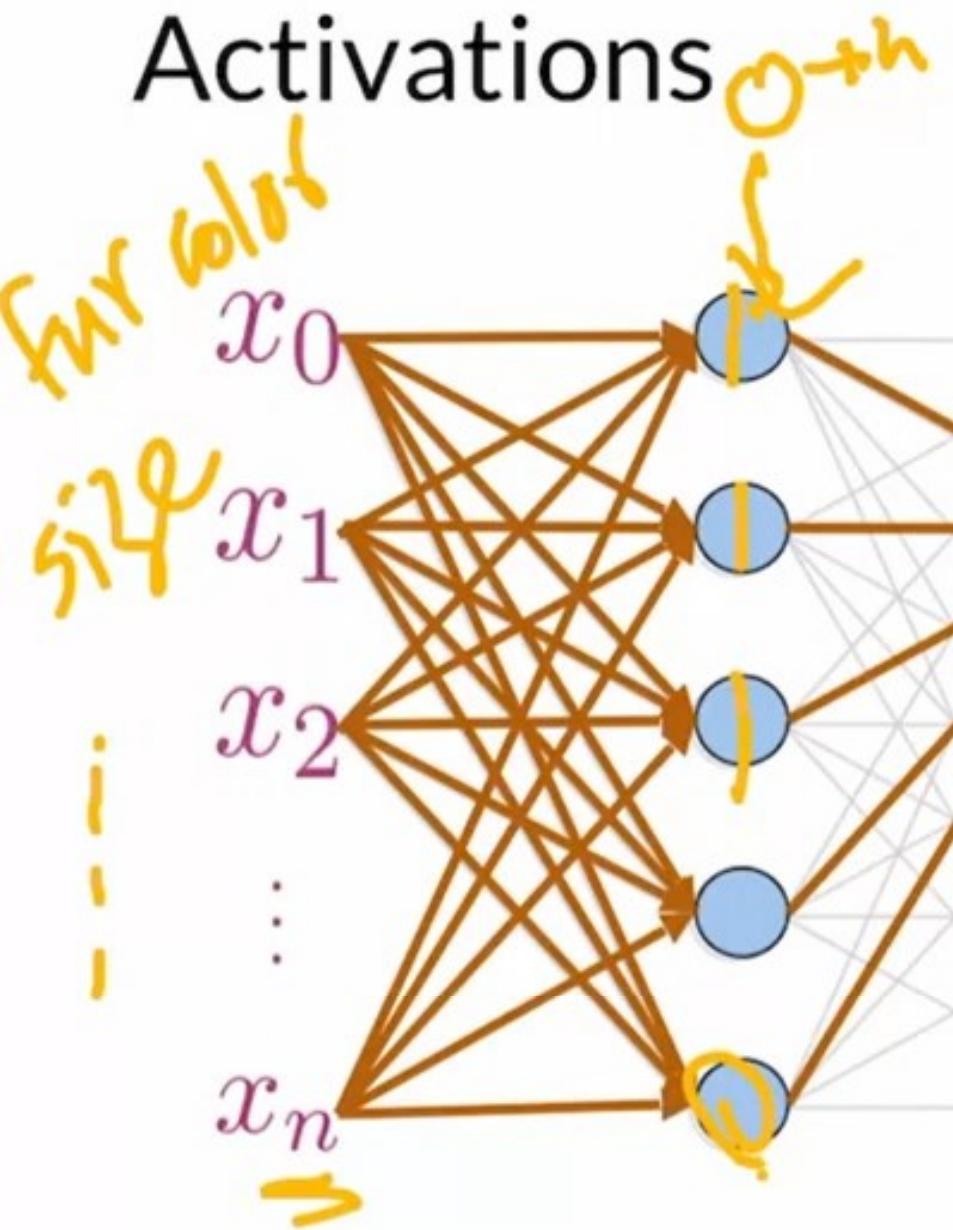
deeplearning.ai

# Activations (Basic Properties)

# Outline

- What are activations
- Reasoning behind non-linear differential activations





$$z_i^{[l]} = \sum_{i=0}^n W_i^{[l]} a_i^{[l-1]} + b$$

$$a_i^{[l]} = g^{[l]}(z_i^{[l]})$$

Differentiable non-linear function

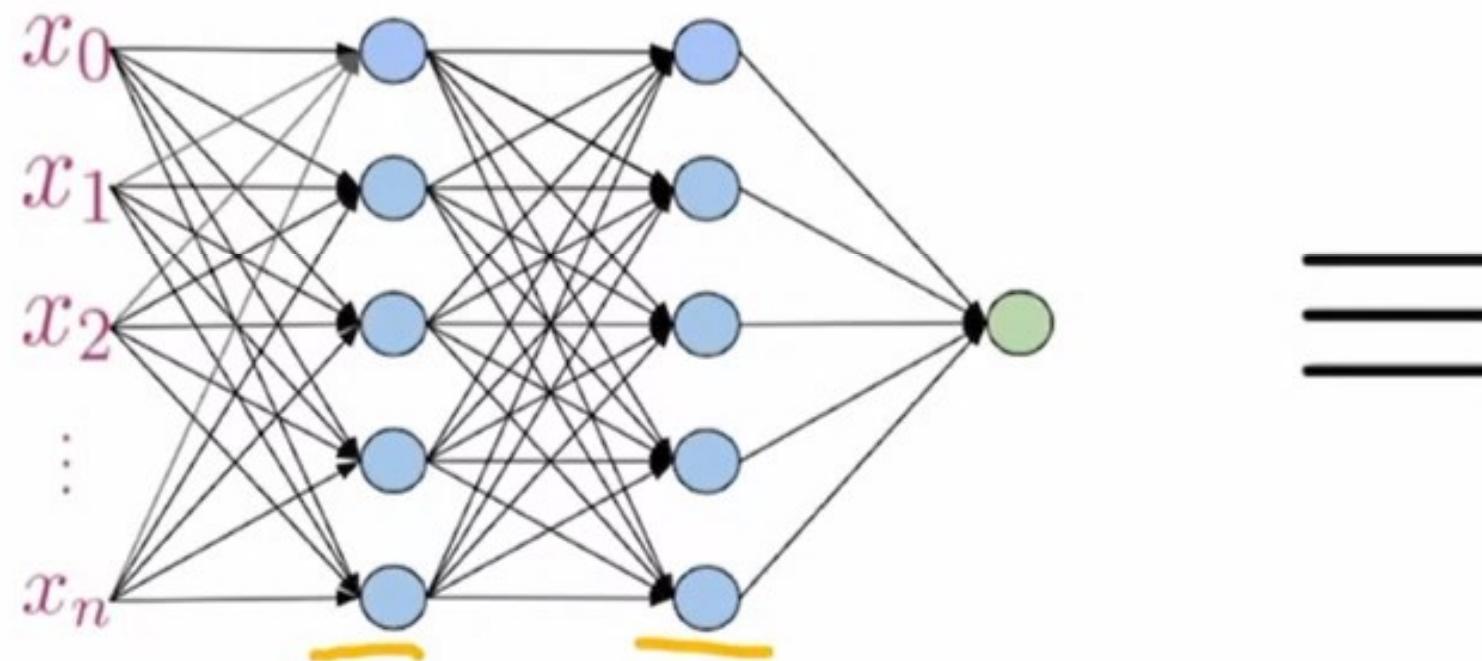
# Activations

$$a_i^{[l]} = g^{[l]}(z_i^{[l]})$$

Differentiable  
non-linear  
function

1. Differentiable for backpropagation

2. Non-linear to compute complex features, **if not**:



$$WX + b$$

Linear  
regression

# Summary

- Activation functions are non-linear and differentiable
- Differentiable for backpropagation
- Non-linear to approximate complex functions

$f\mathbf{x}$



deeplearning.ai

# Common Activation Functions

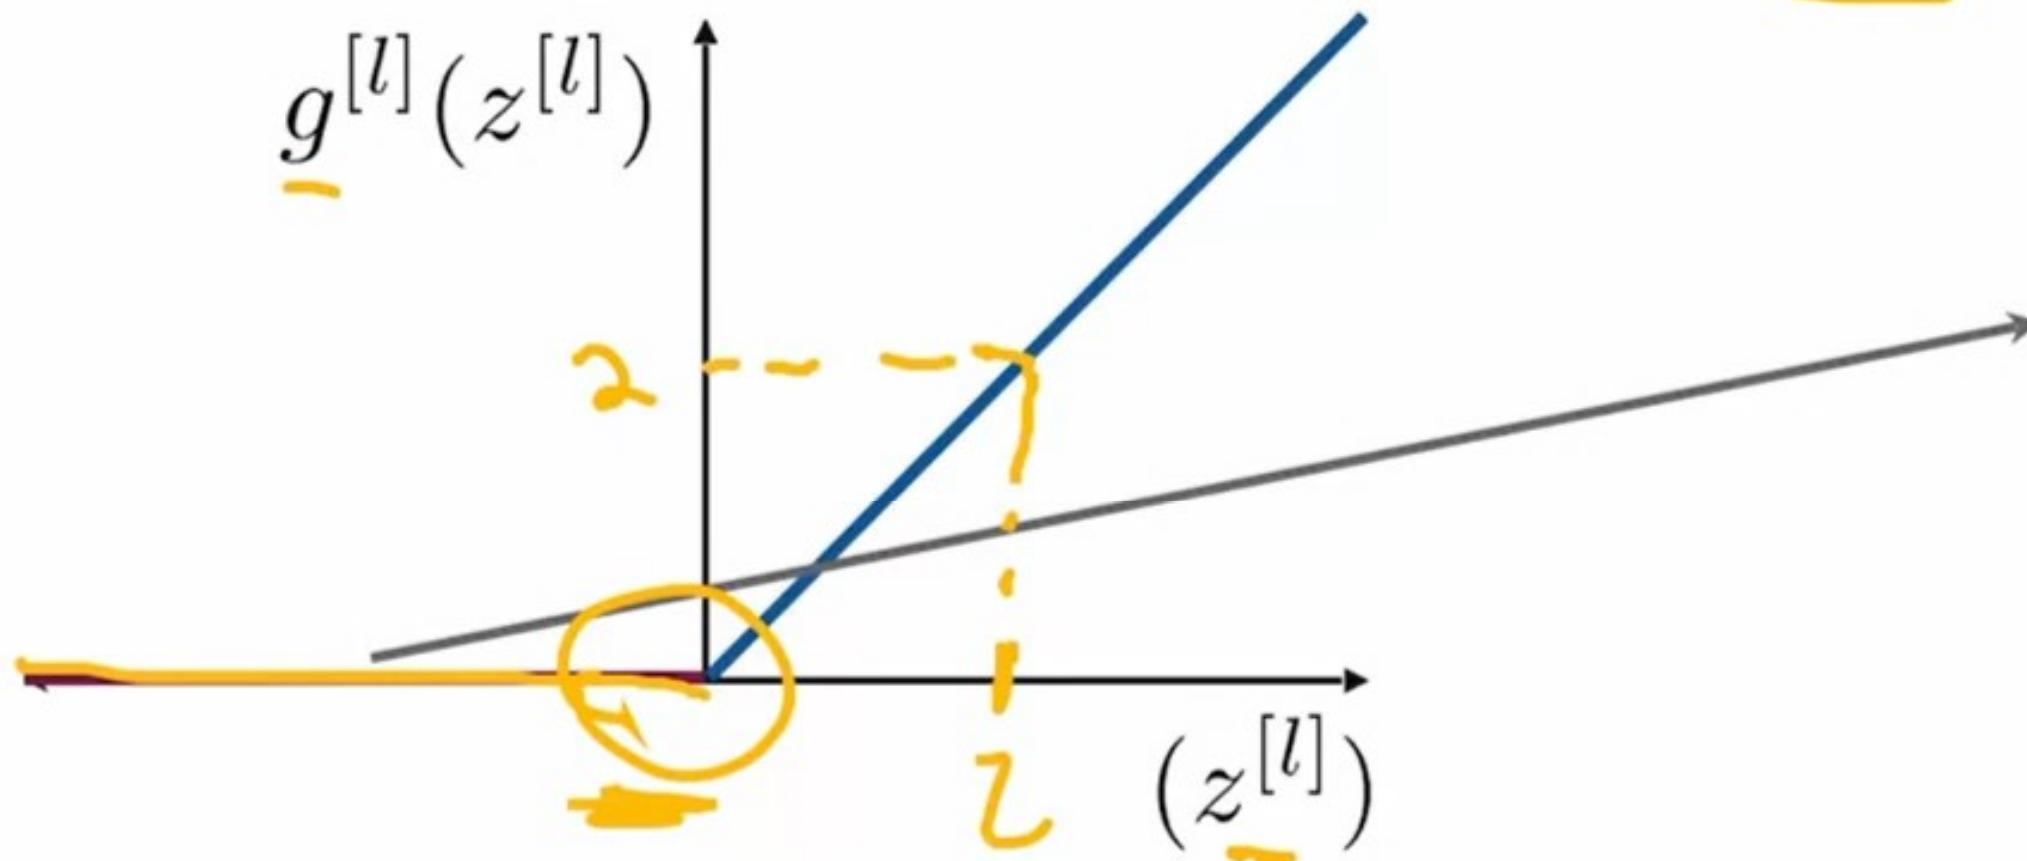
# Outline

- Common activations and their structure
  - ReLU
  - Leaky ReLU
  - Sigmoid
  - Tanh



## Activations: ReLU

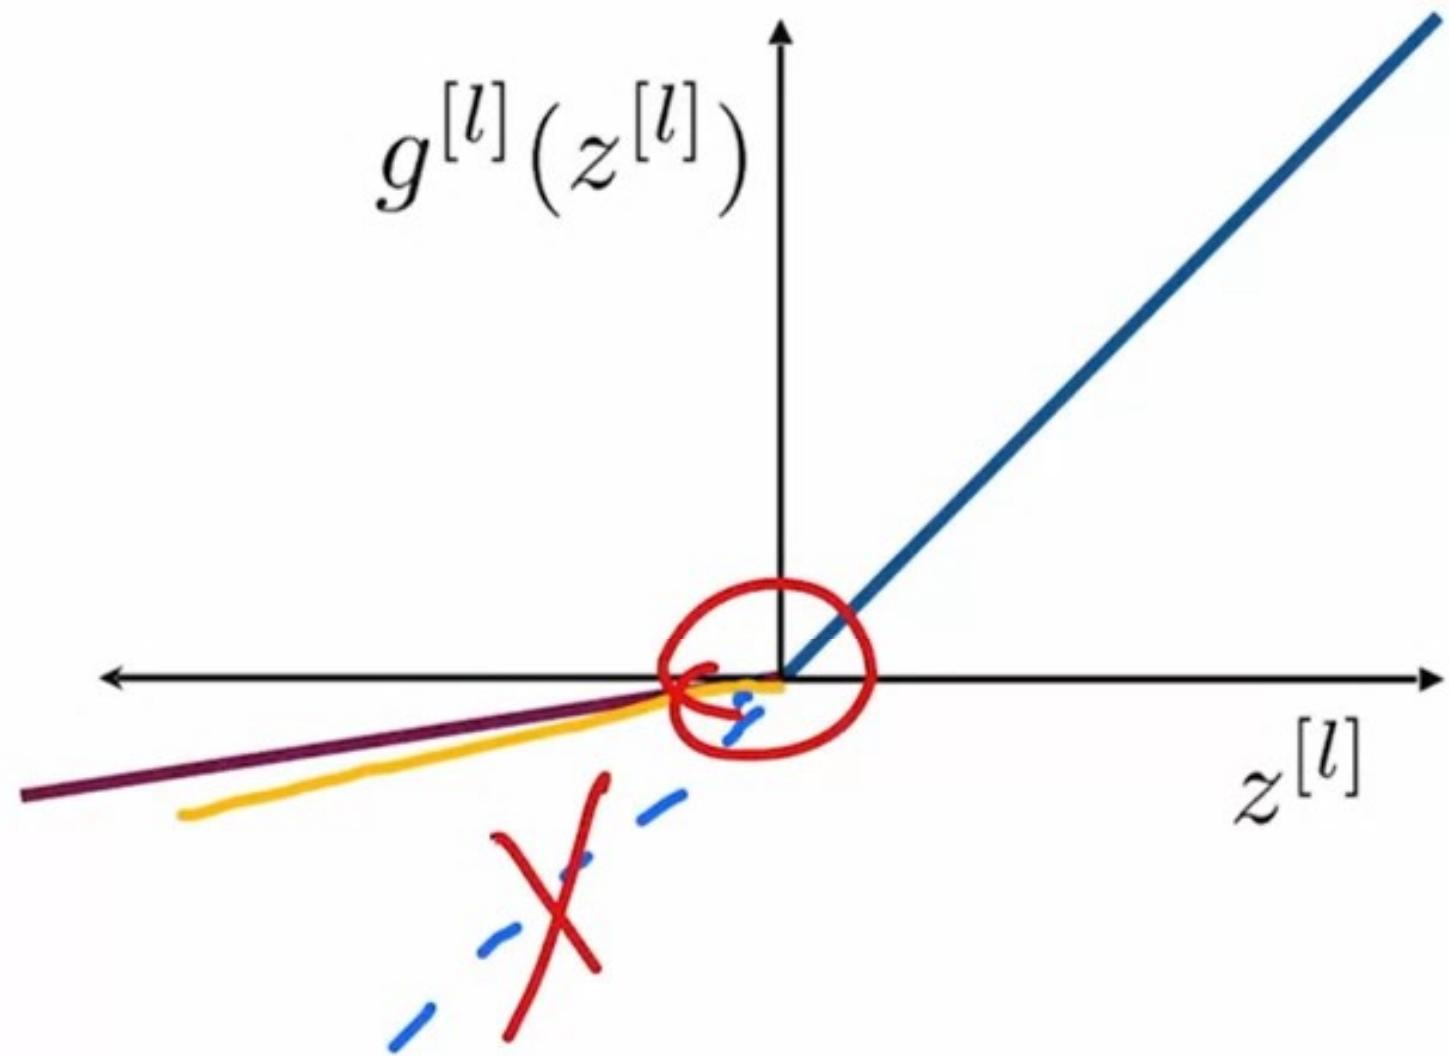
ReLU = Rectified Linear Unit



ReLU  
 $\downarrow$   
 $g^{[l]}(z^{[l]}) = \max(0, z^{[l]})$

Dying ReLU  
problem

## Activations: Leaky ReLU



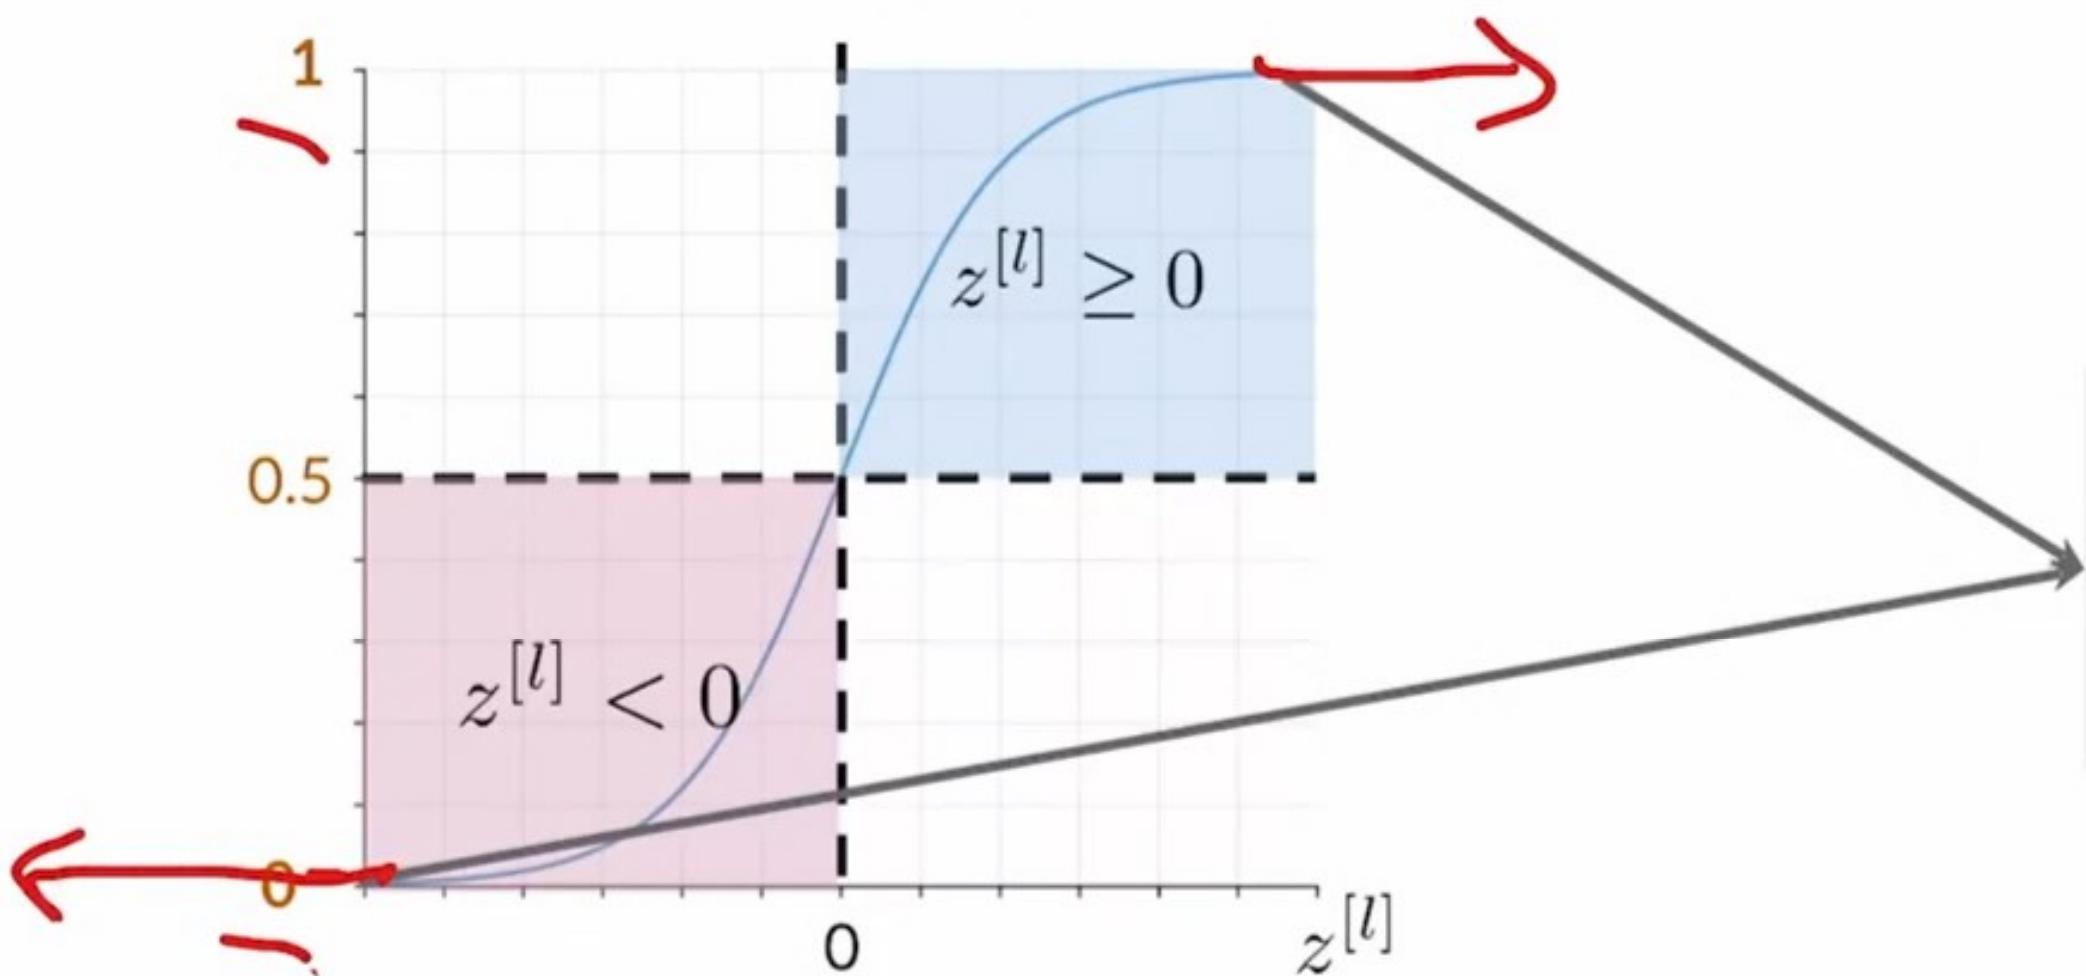
$$g^{[l]}(z^{[l]}) = \max(\underline{az}^{[l]}, \underline{z}^{[l]})$$

Solves the dying  
ReLU problem

0.1

# Activations: Sigmoid

$$g^{[l]}(z^{[l]}) = \frac{1}{1 + e^{-z^{[l]}}}$$

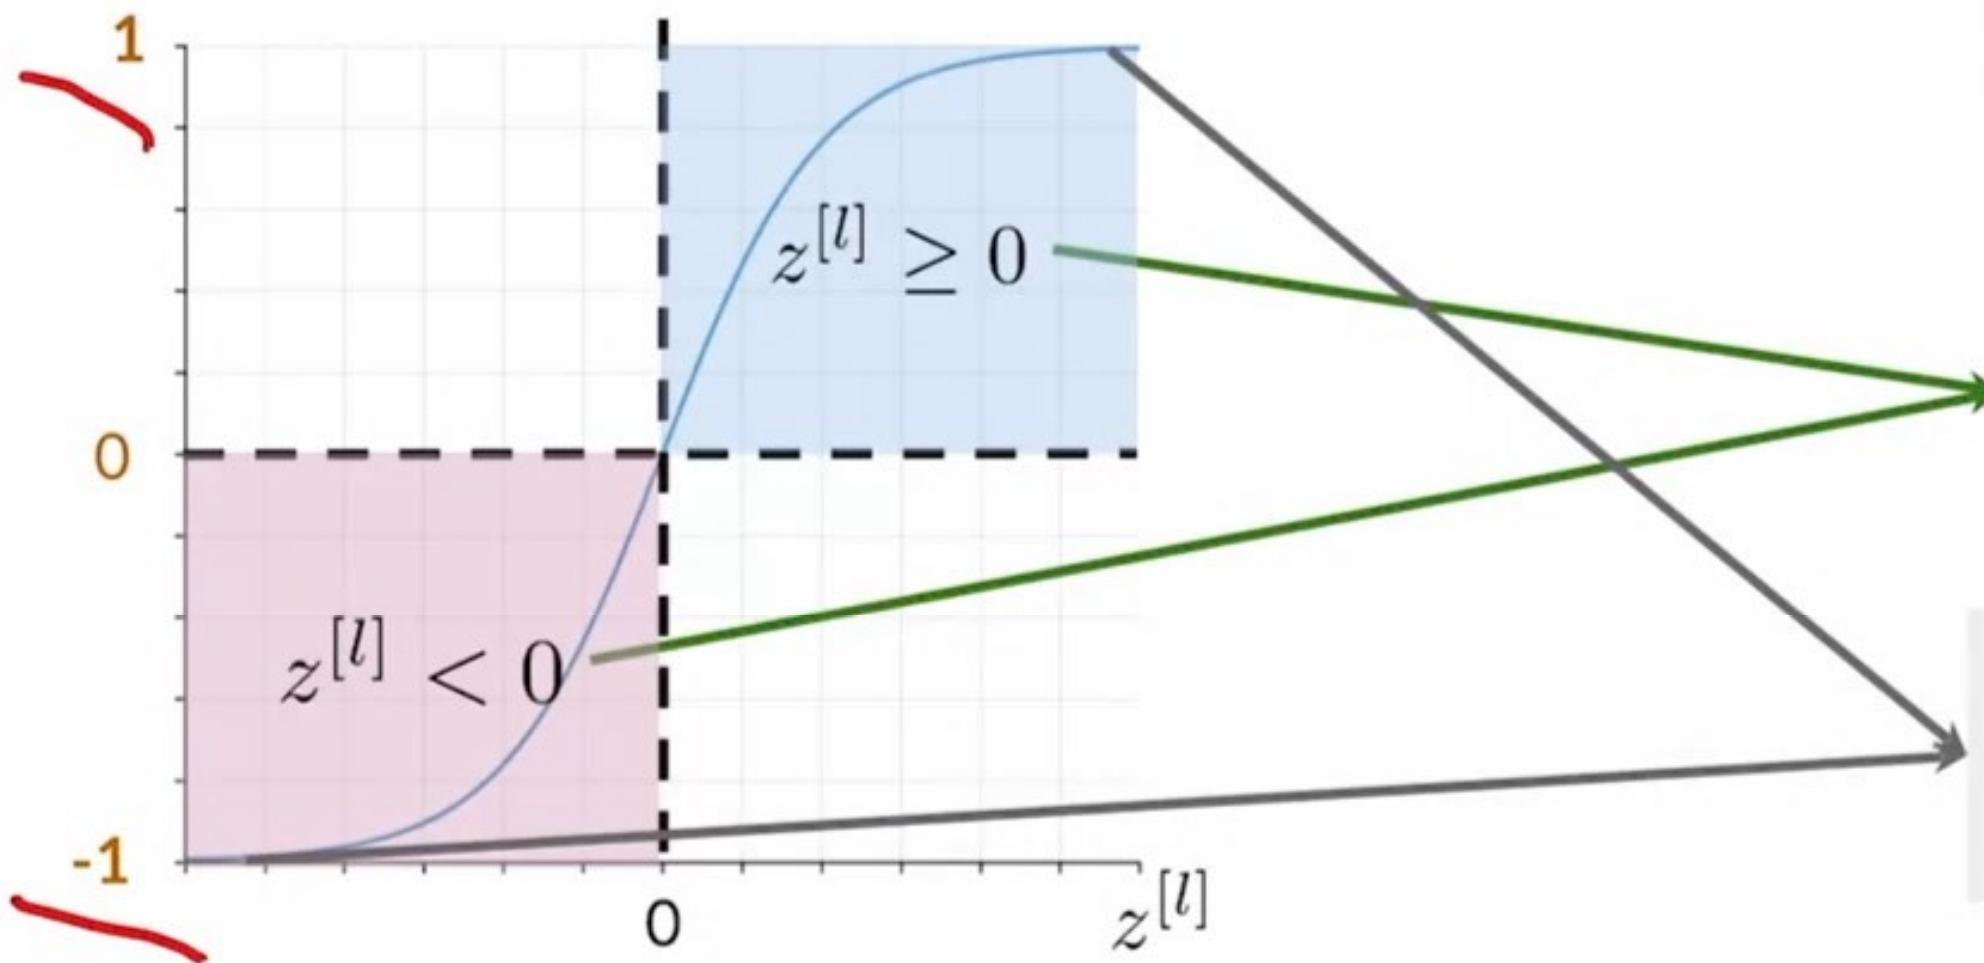


Values between 0 and 1

Vanishing gradient and saturation problems

# Activations: Tanh

$$g^{[l]}(z^{[l]}) = \tanh(z^{[l]})$$



Values between -1 and 1

Keeps the sign of the input

Same issues as Sigmoid

# Summary

- ReLU activations suffer from dying ReLU
- Leaky ReLU solve the dying ReLU problem
- Sigmoid and Tanh have vanishing gradient and saturation problems





deeplearning.ai

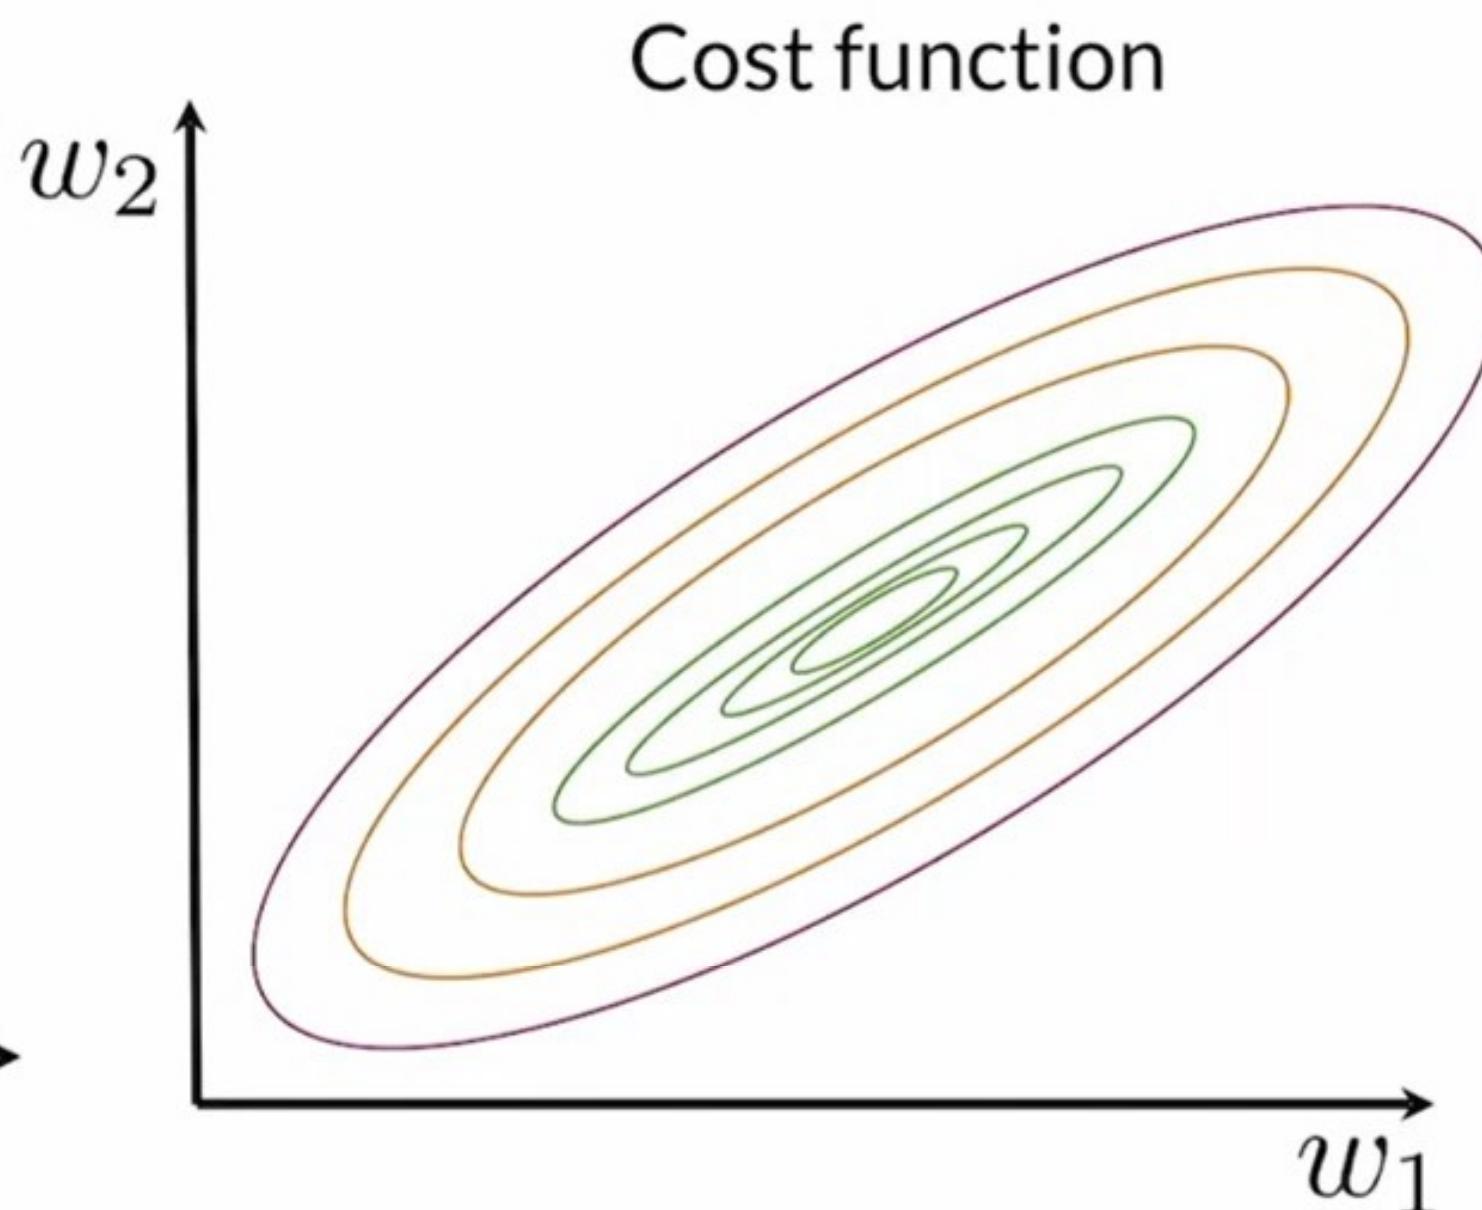
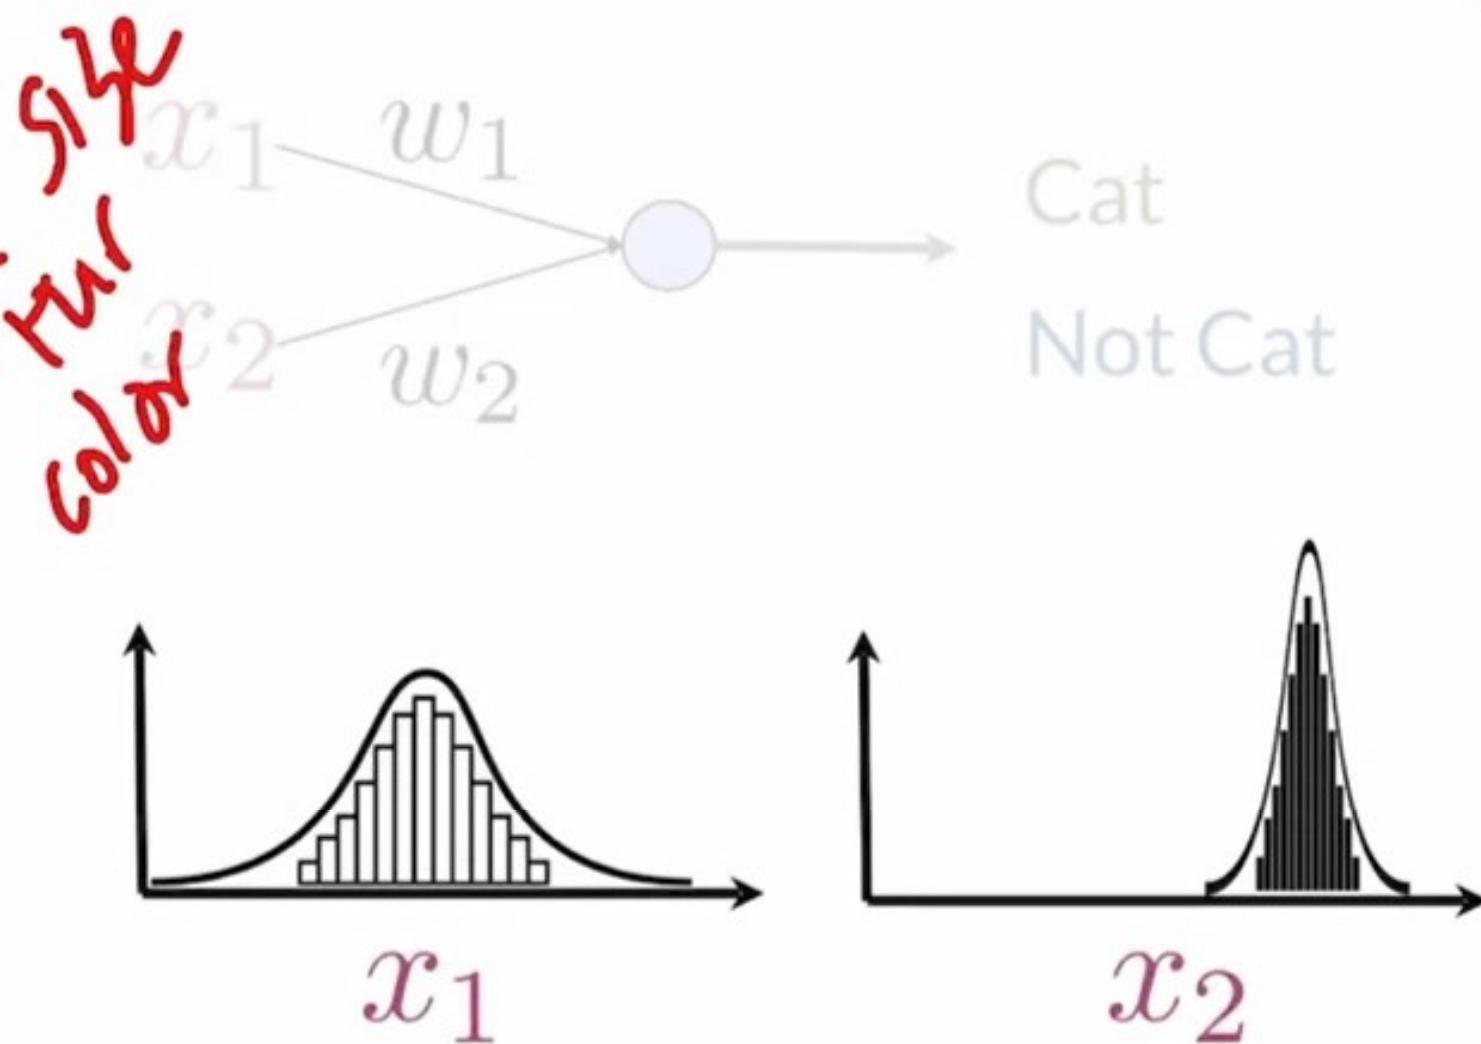
# Batch Normalization (Explained)

# Outline

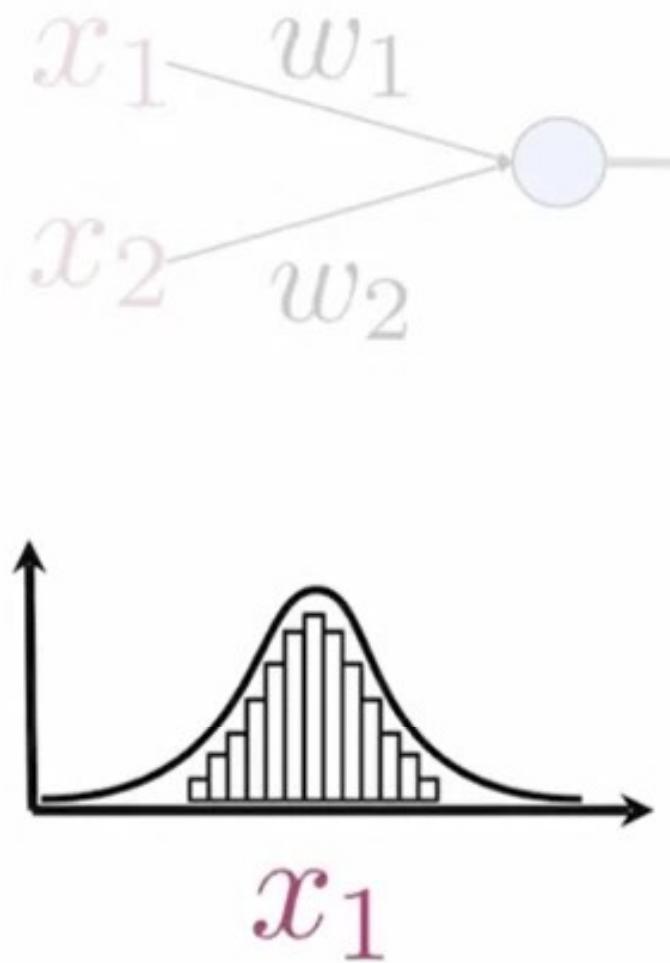
- How normalization helps models
- Internal covariate shift
- Batch normalization



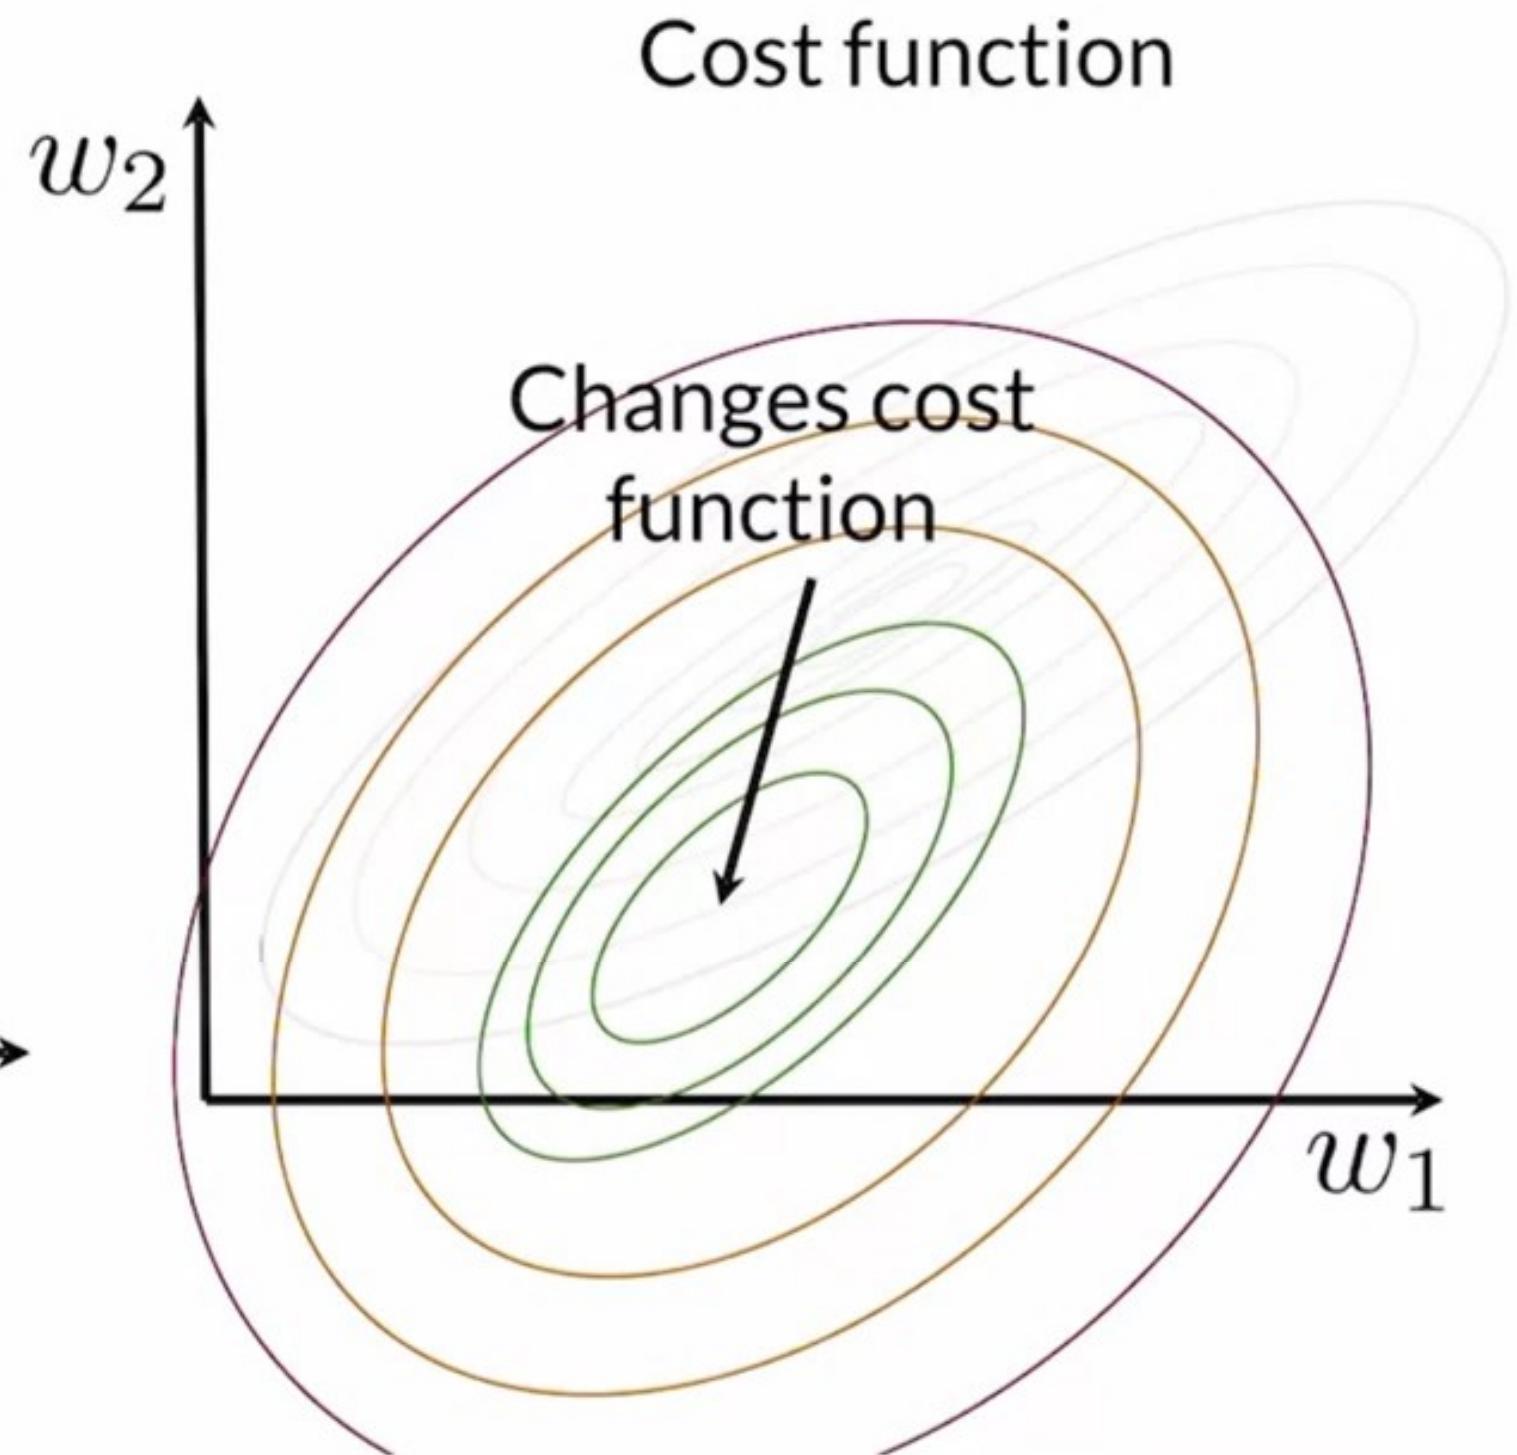
# Different Distributions



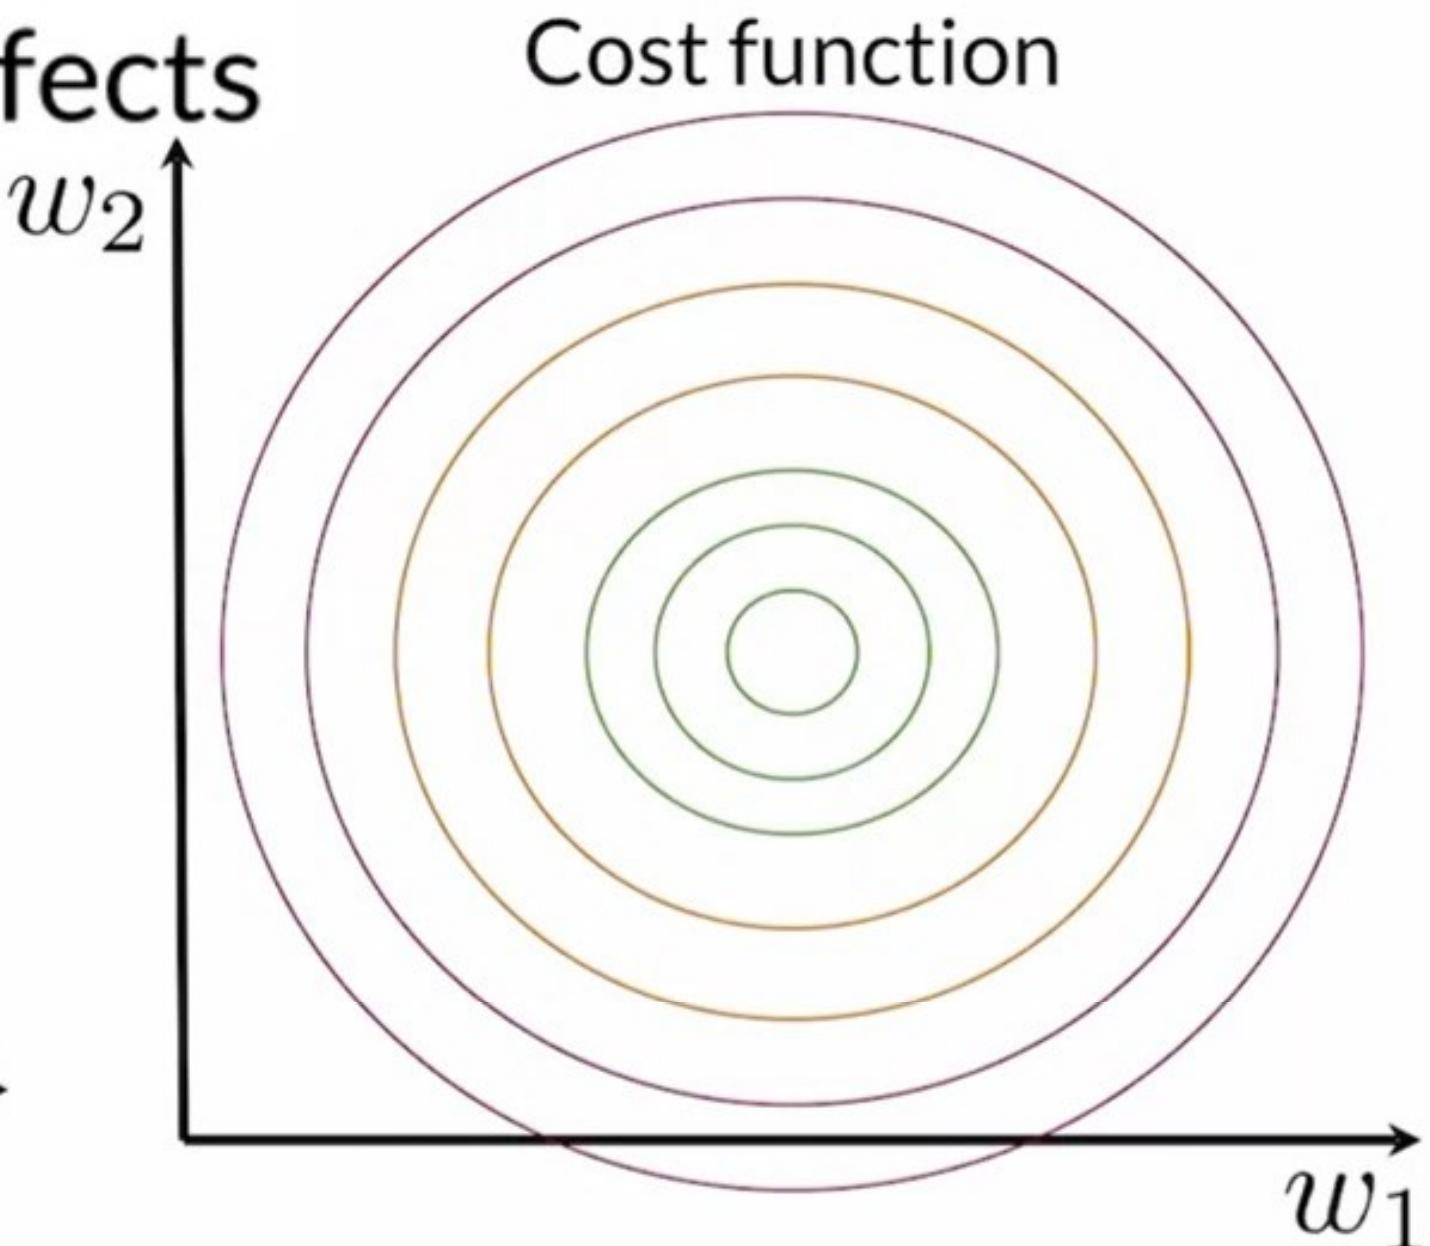
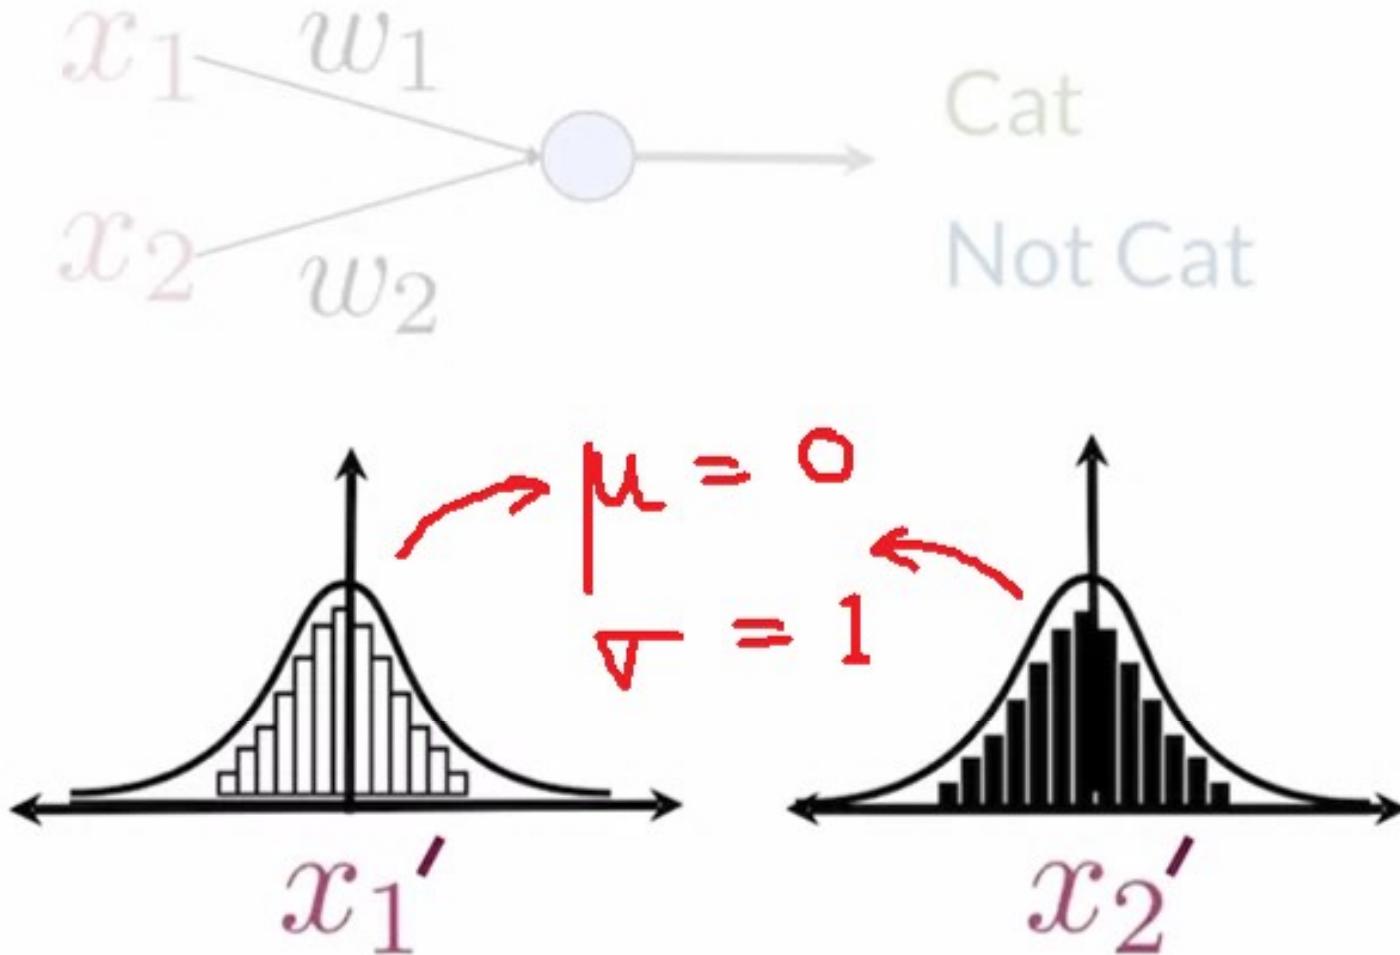
# Covariate Shift



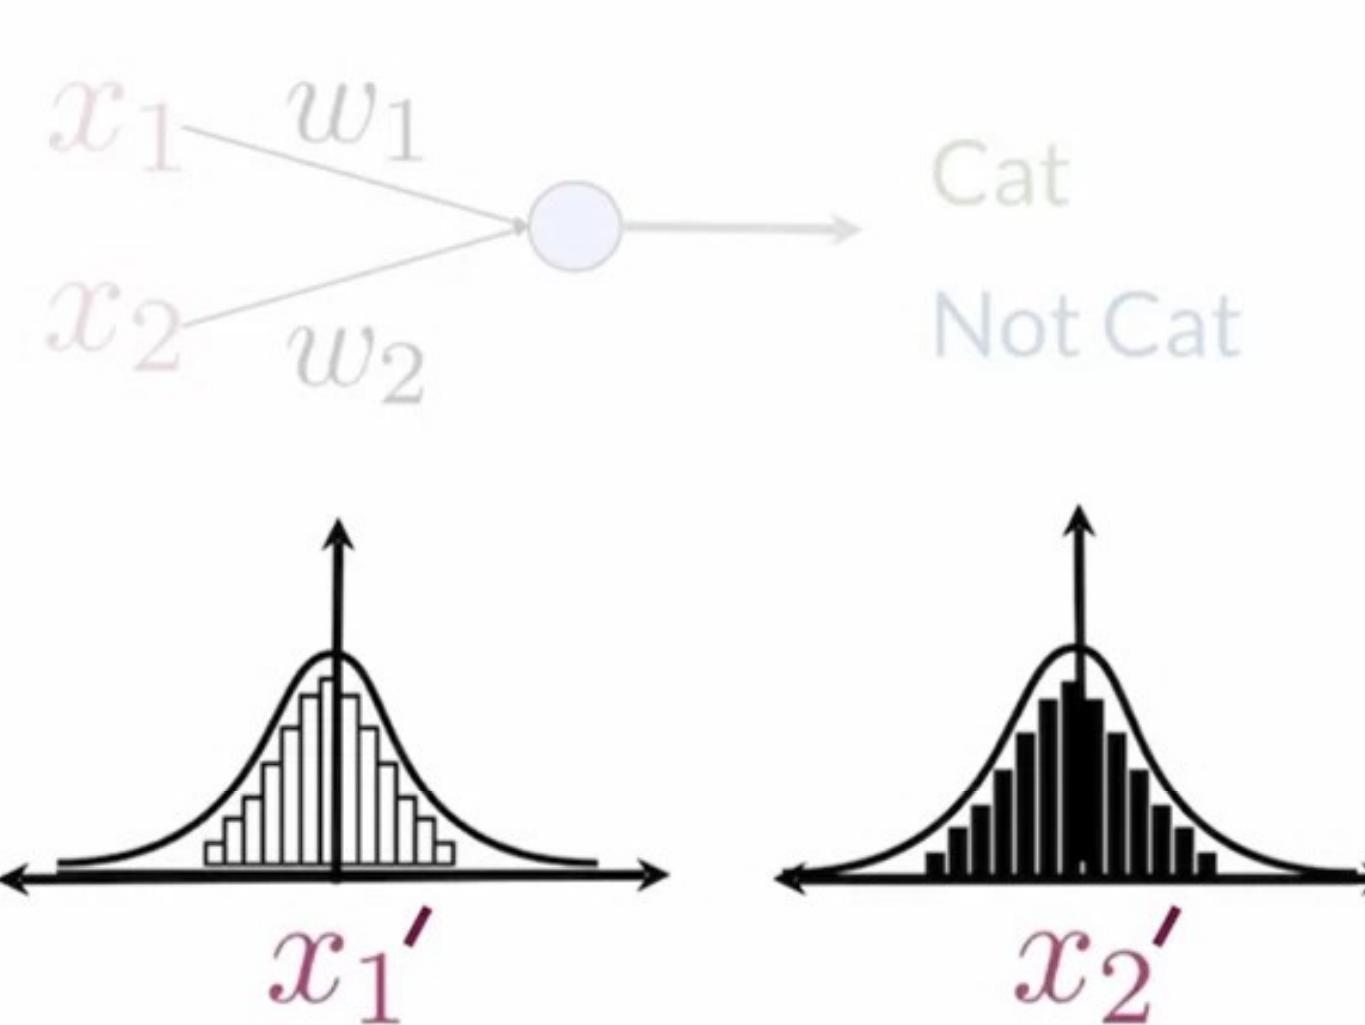
Cat  
Not Cat  
Change in data distribution



# Normalization and Its Effects



# Normalization and Its Effects

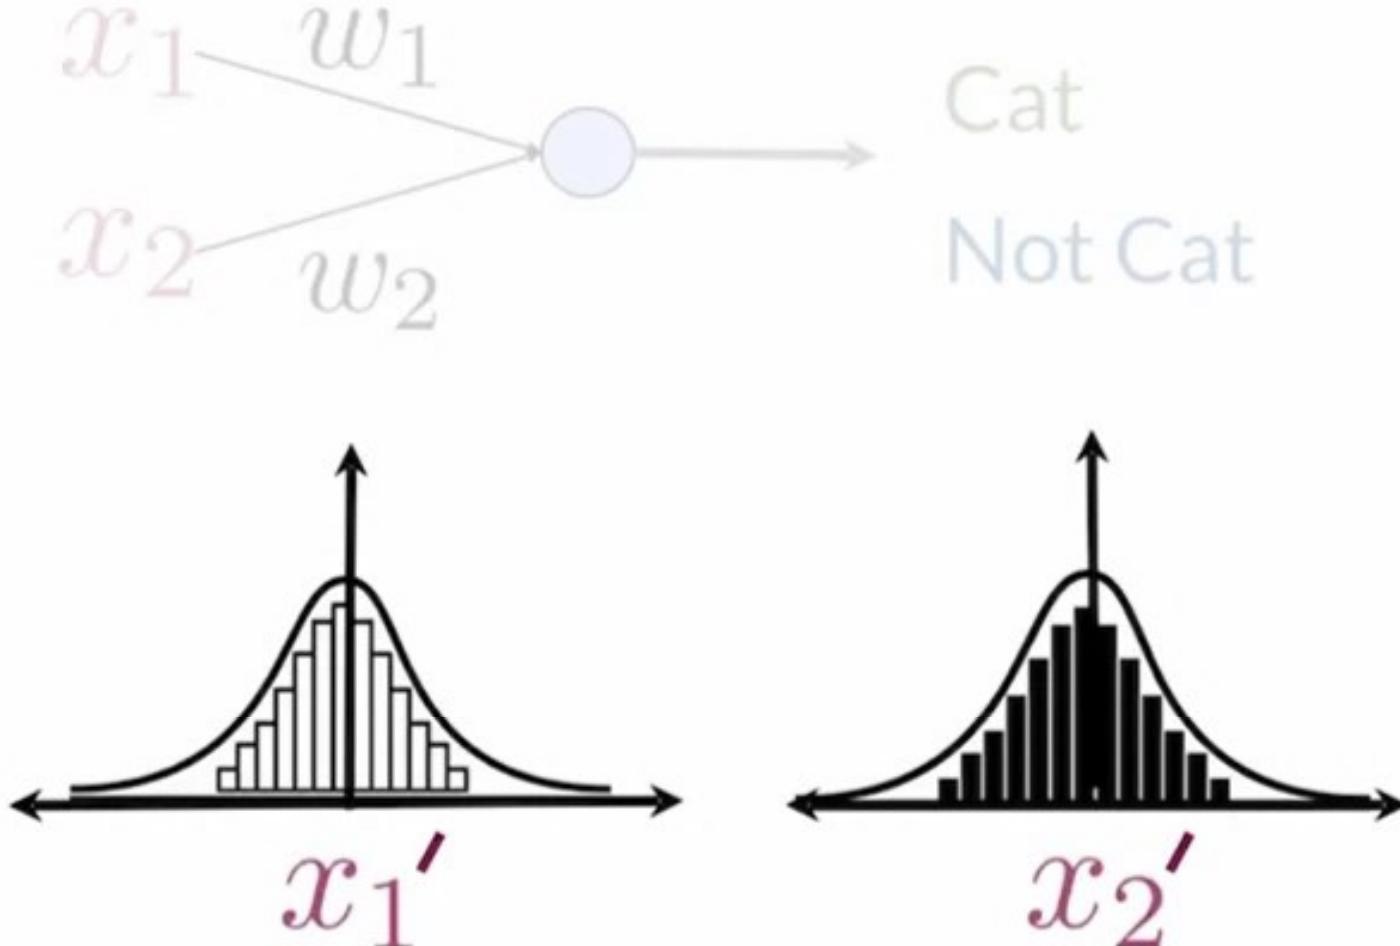


**Around** mean at 0  
and std. at 1

Training data uses  
batch stats

Test data uses  
training stats

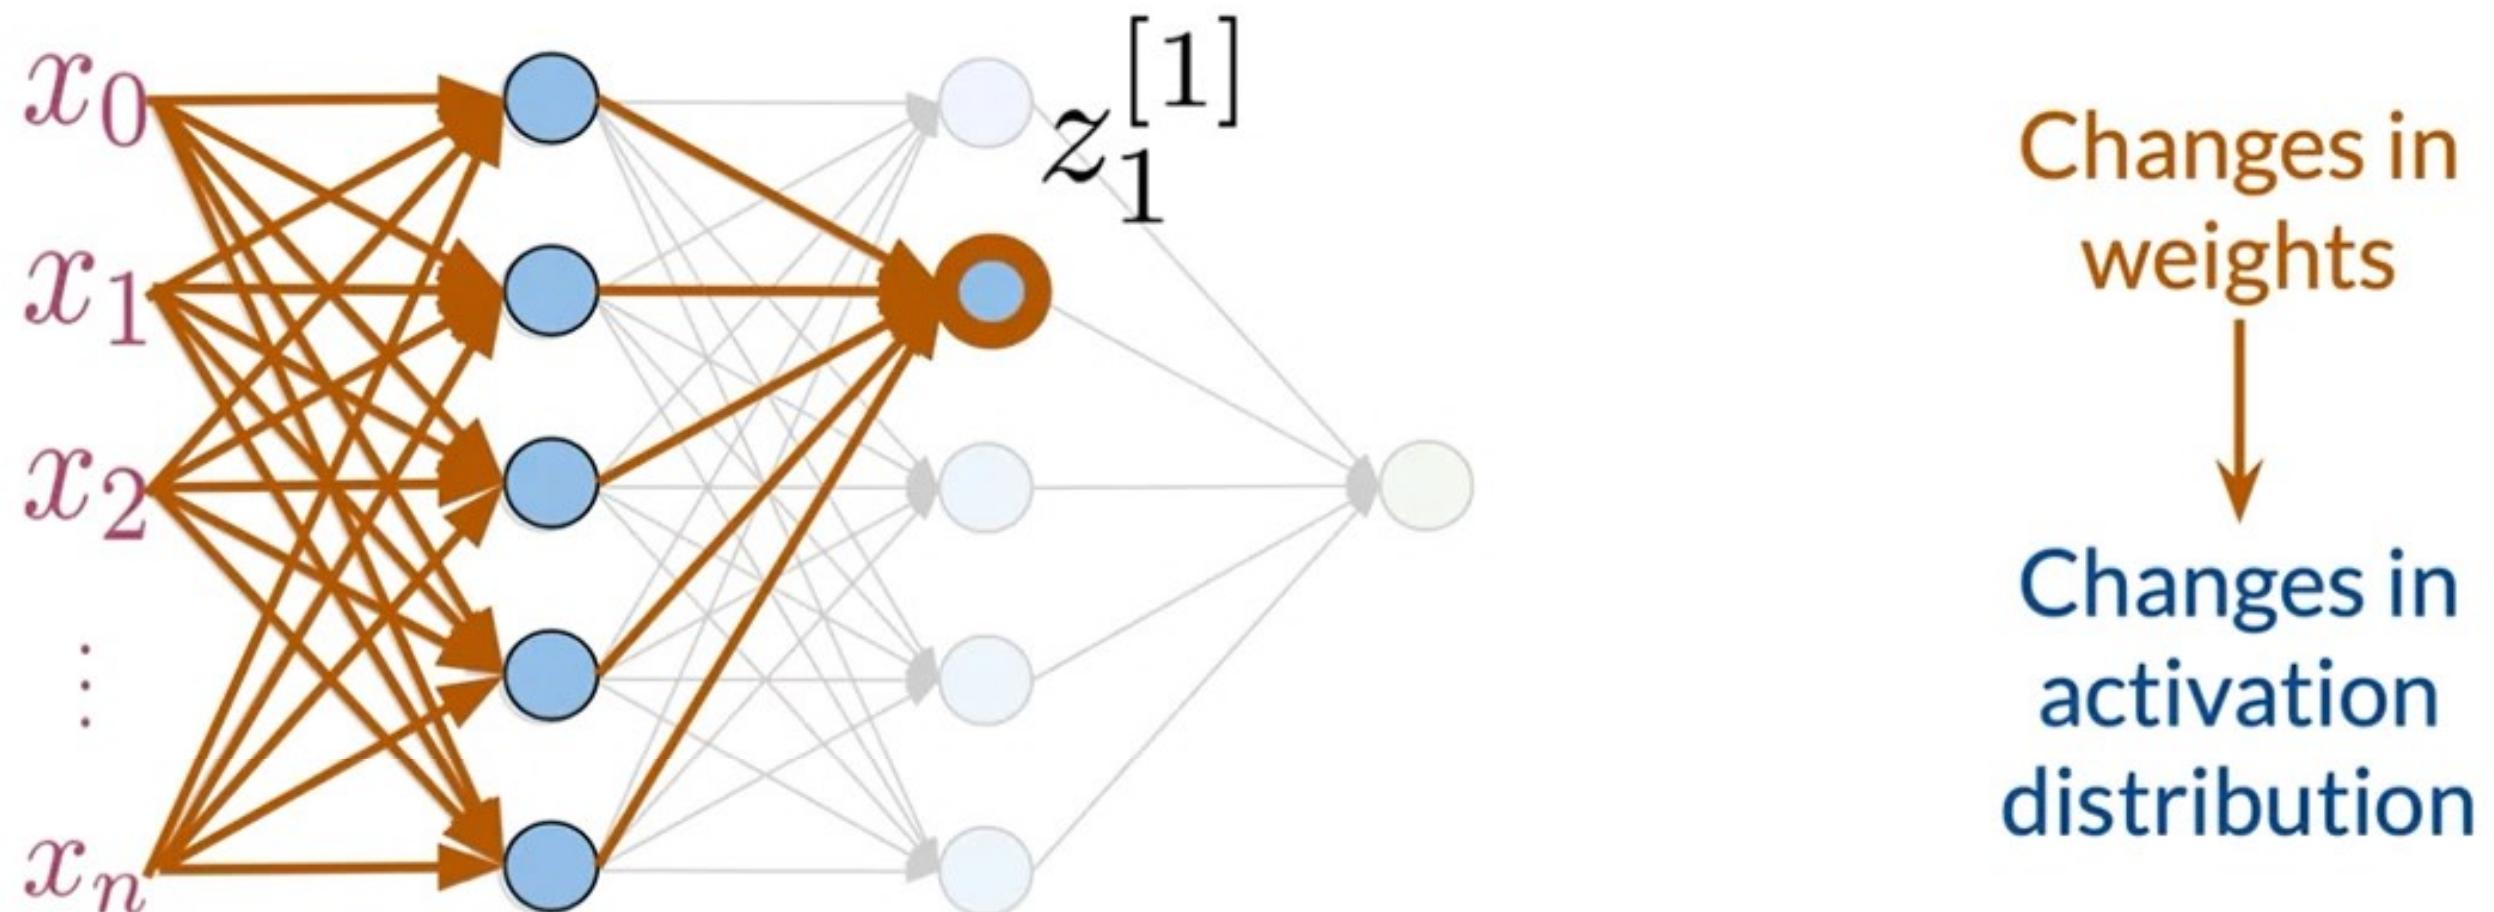
# Normalization and Its Effects



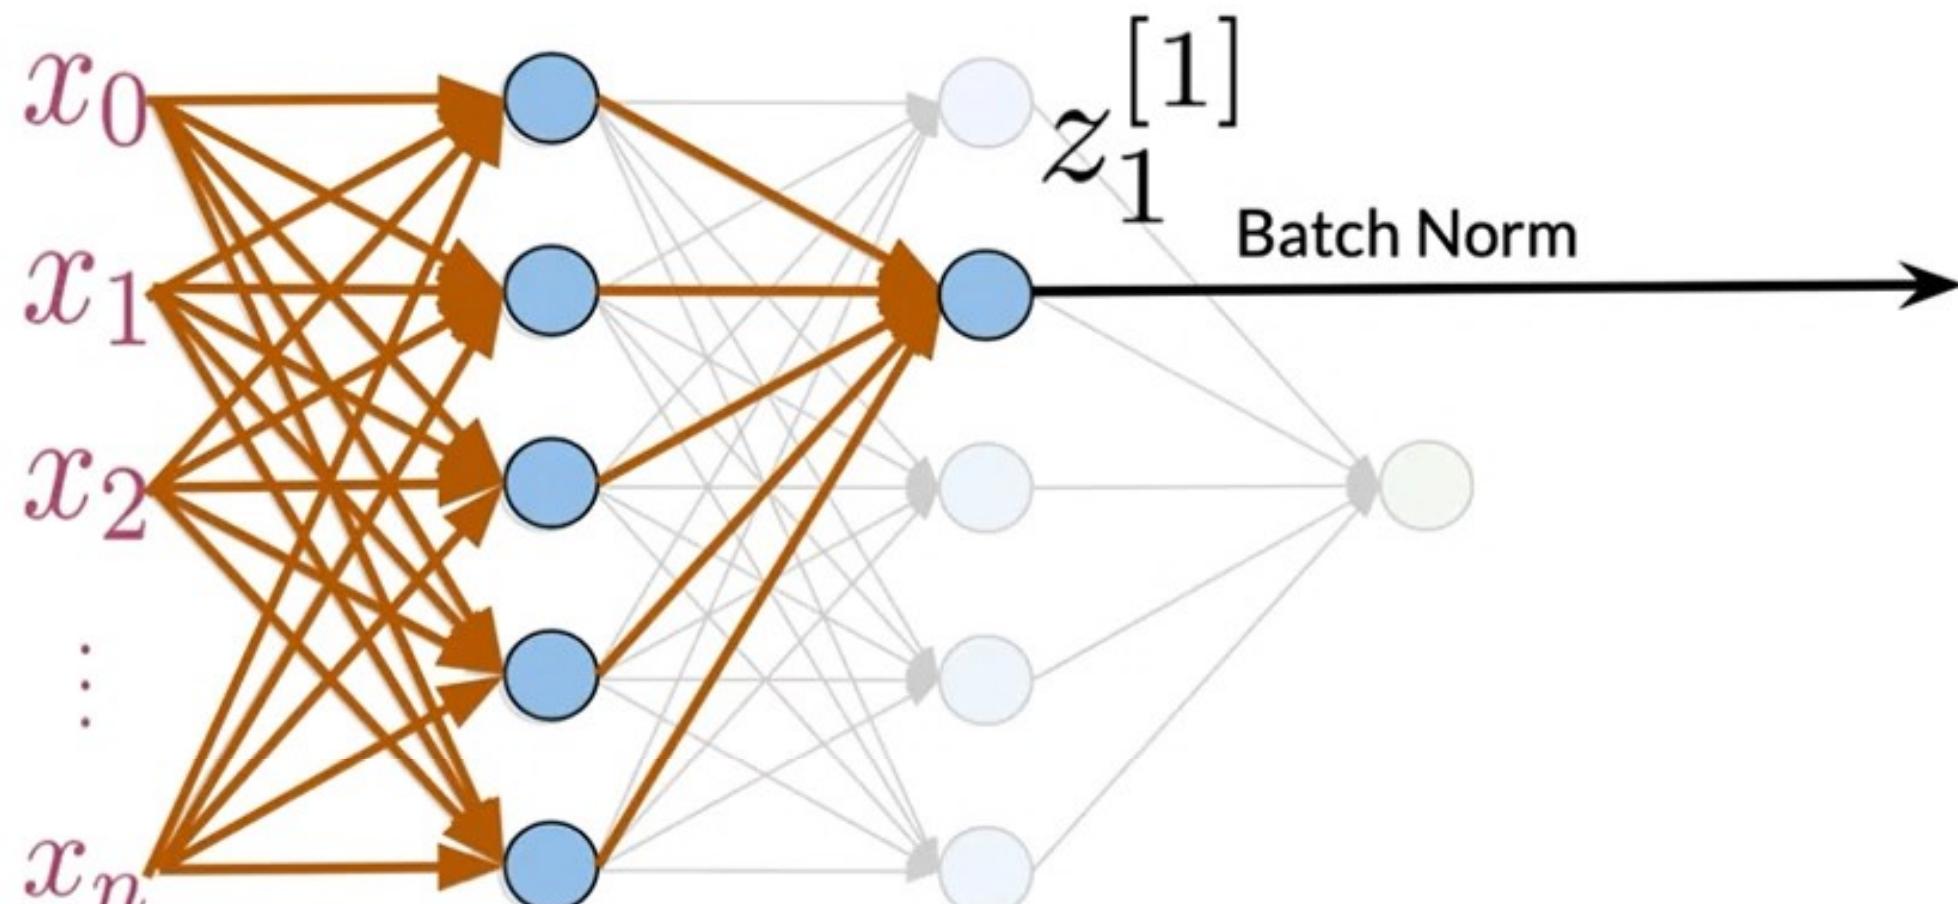
**Around** mean at 0  
and std. at 1

Reduction of  
covariate shift

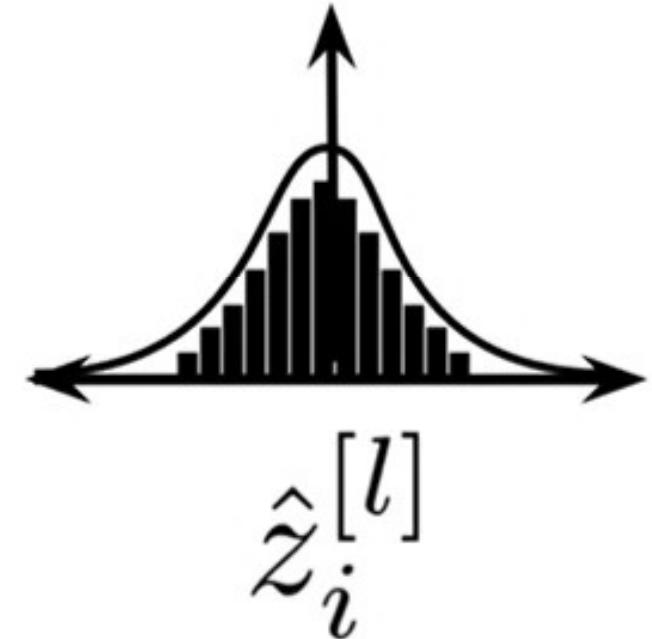
# Internal Covariate Shift



# Batch Normalization



Normalizes the  
input for each  
neuron



# Summary

- Batch normalization smooths the cost function
- Batch normalization reduces the internal covariate shift
- Batch normalization speeds up learning!



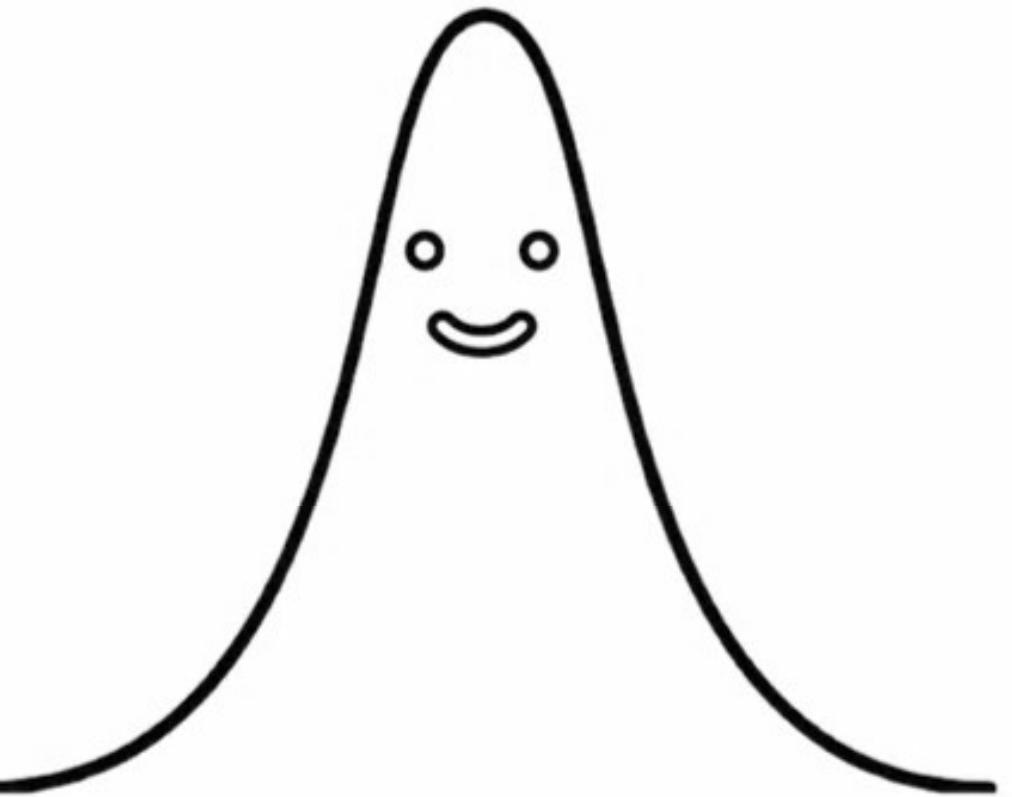


deeplearning.ai

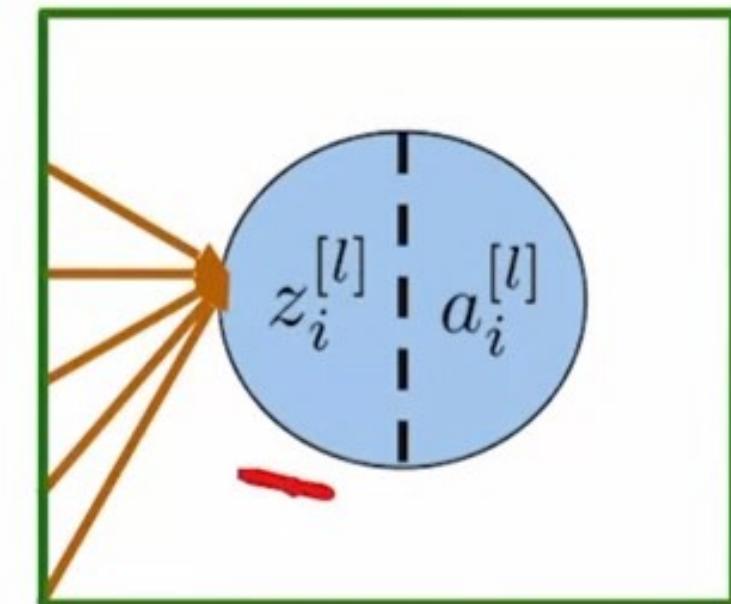
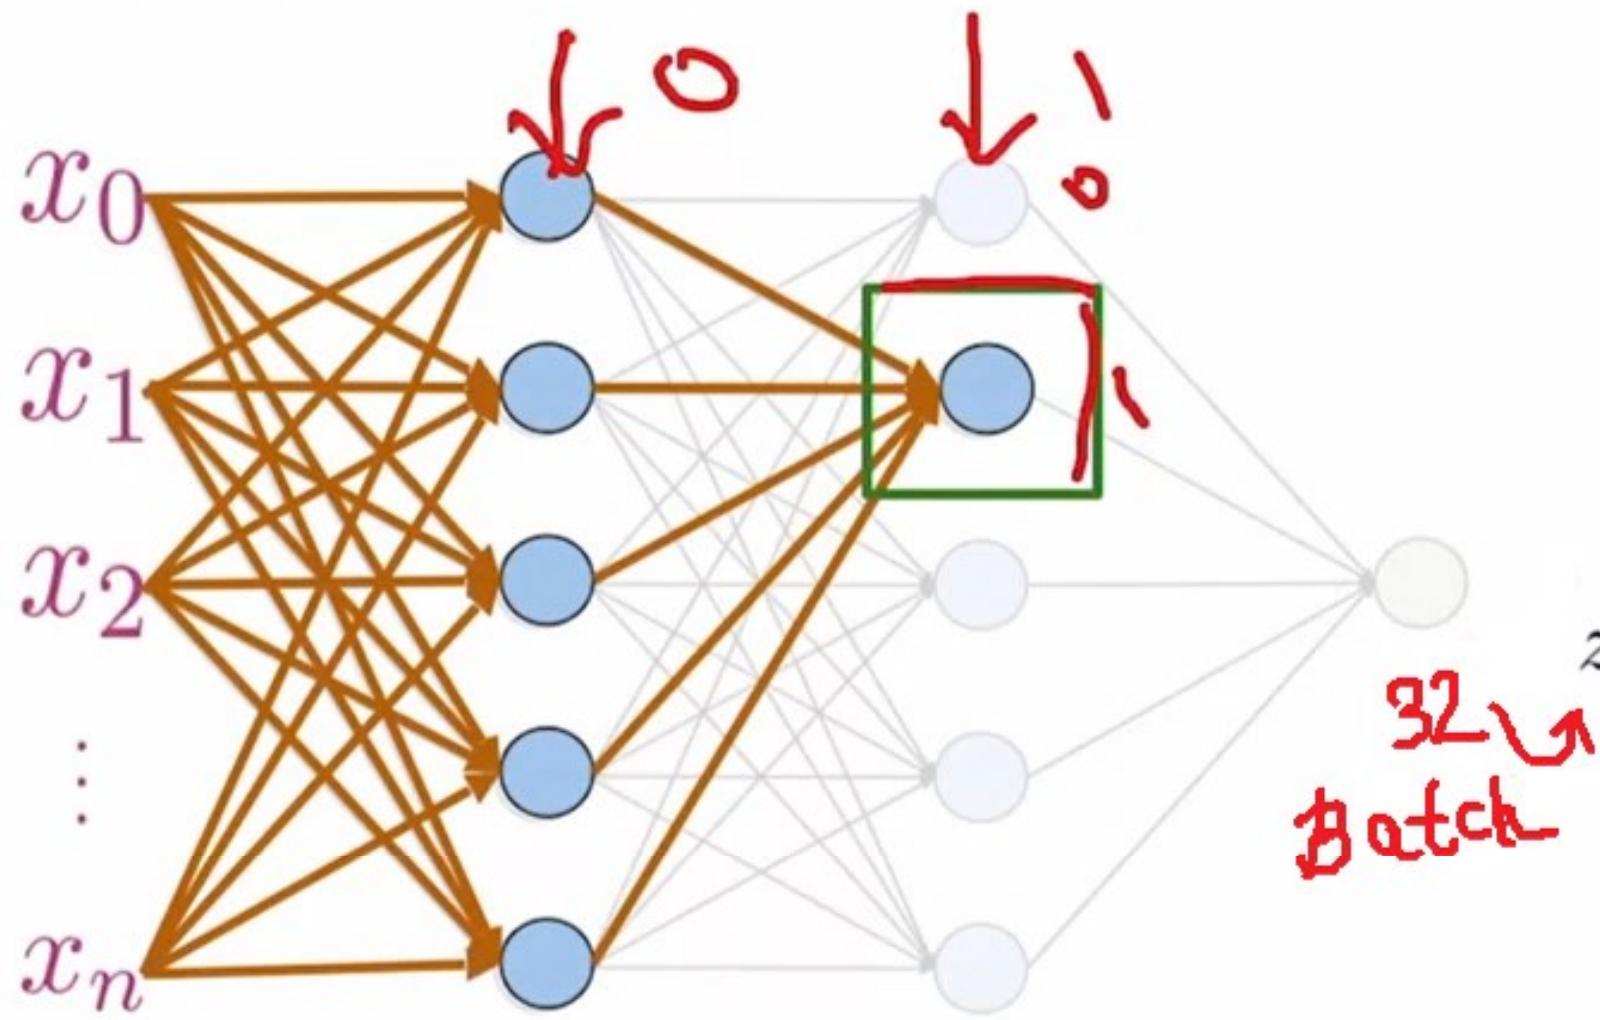
# Batch Normalization (Procedure)

# Outline

- Batch norm for training
- Batch norm for testing



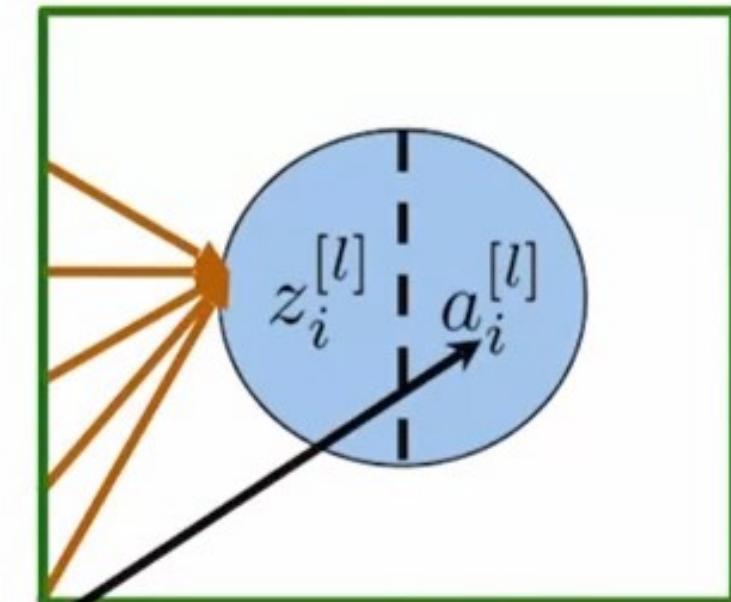
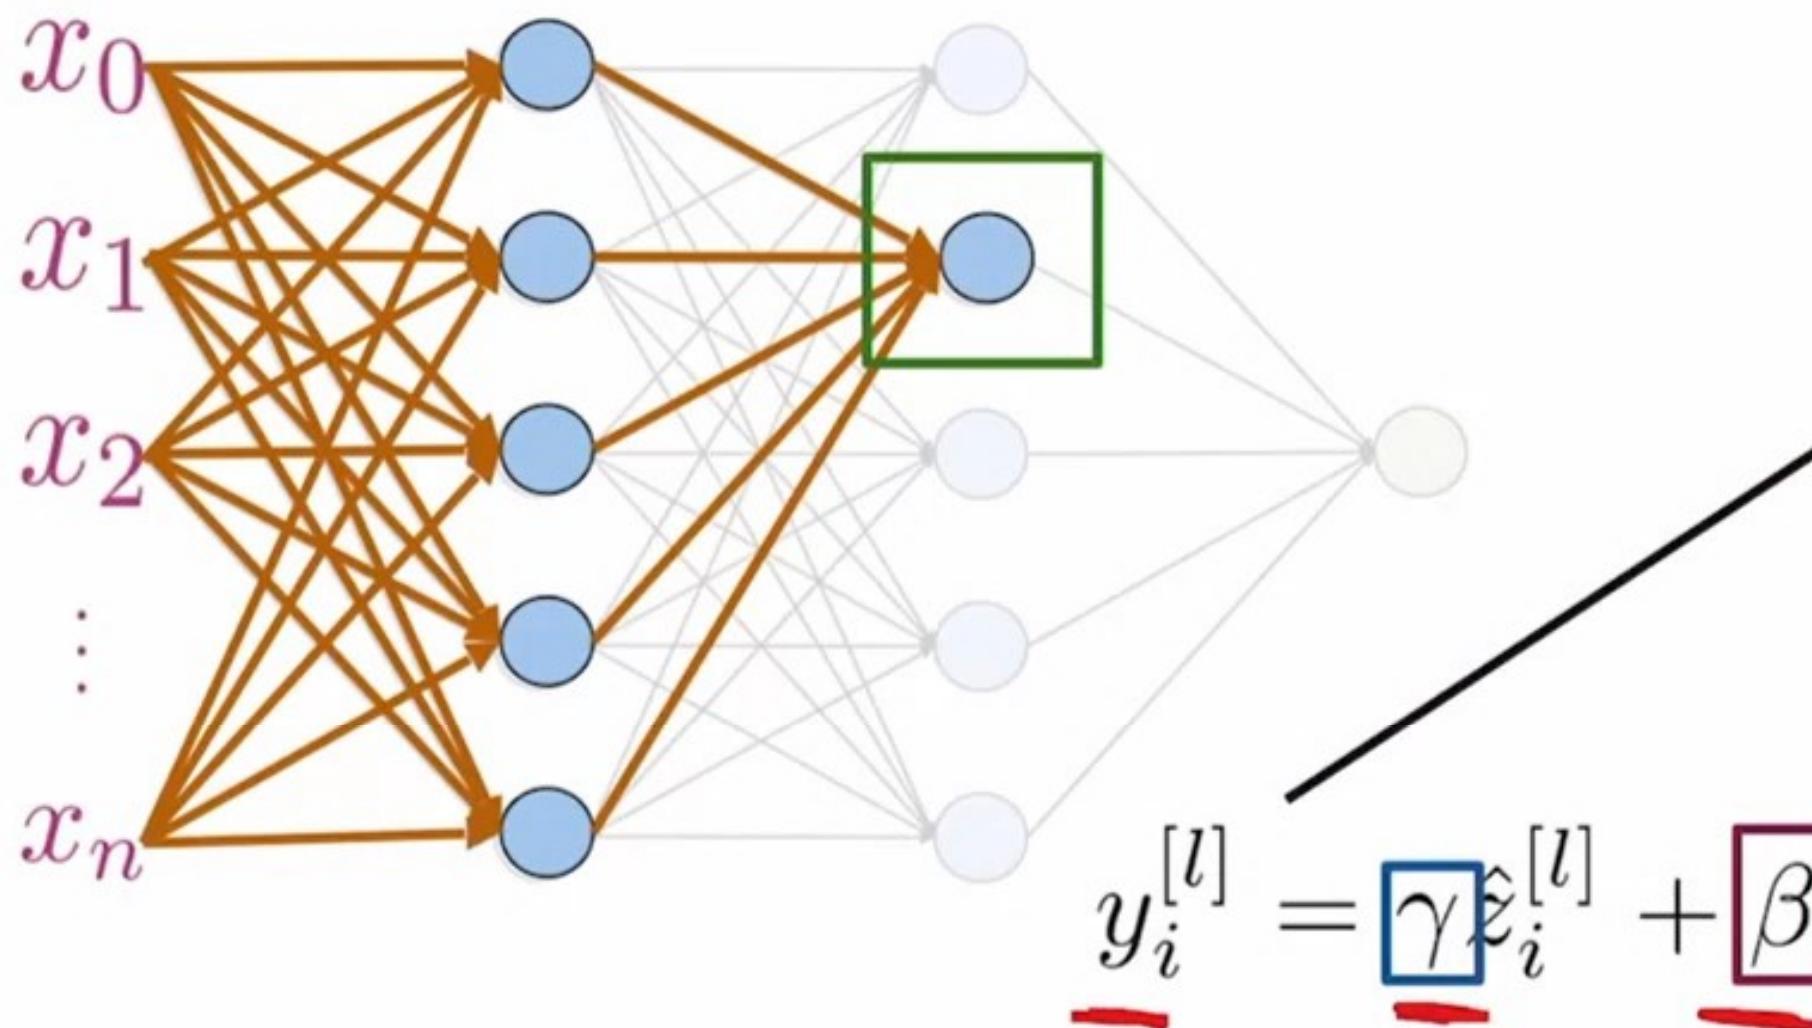
# Batch Normalization: Training



$$z_i^{[l]} = \sum_{i=0} W_i^{[l]} a_i^{[l-1]} \longrightarrow \text{For every example in the batch}$$

$$\hat{z}_i^{[l]} = \frac{z_i^{[l]} - \mu_{z_i^{[l]}}}{\sqrt{\sigma_{z_i^{[l]}}^2 + \epsilon}} \quad \begin{array}{l} \text{Batch mean of } z_i^{[l]} \\ \text{Batch std of } z_i^{[l]} \end{array}$$

# Batch Normalization: Training

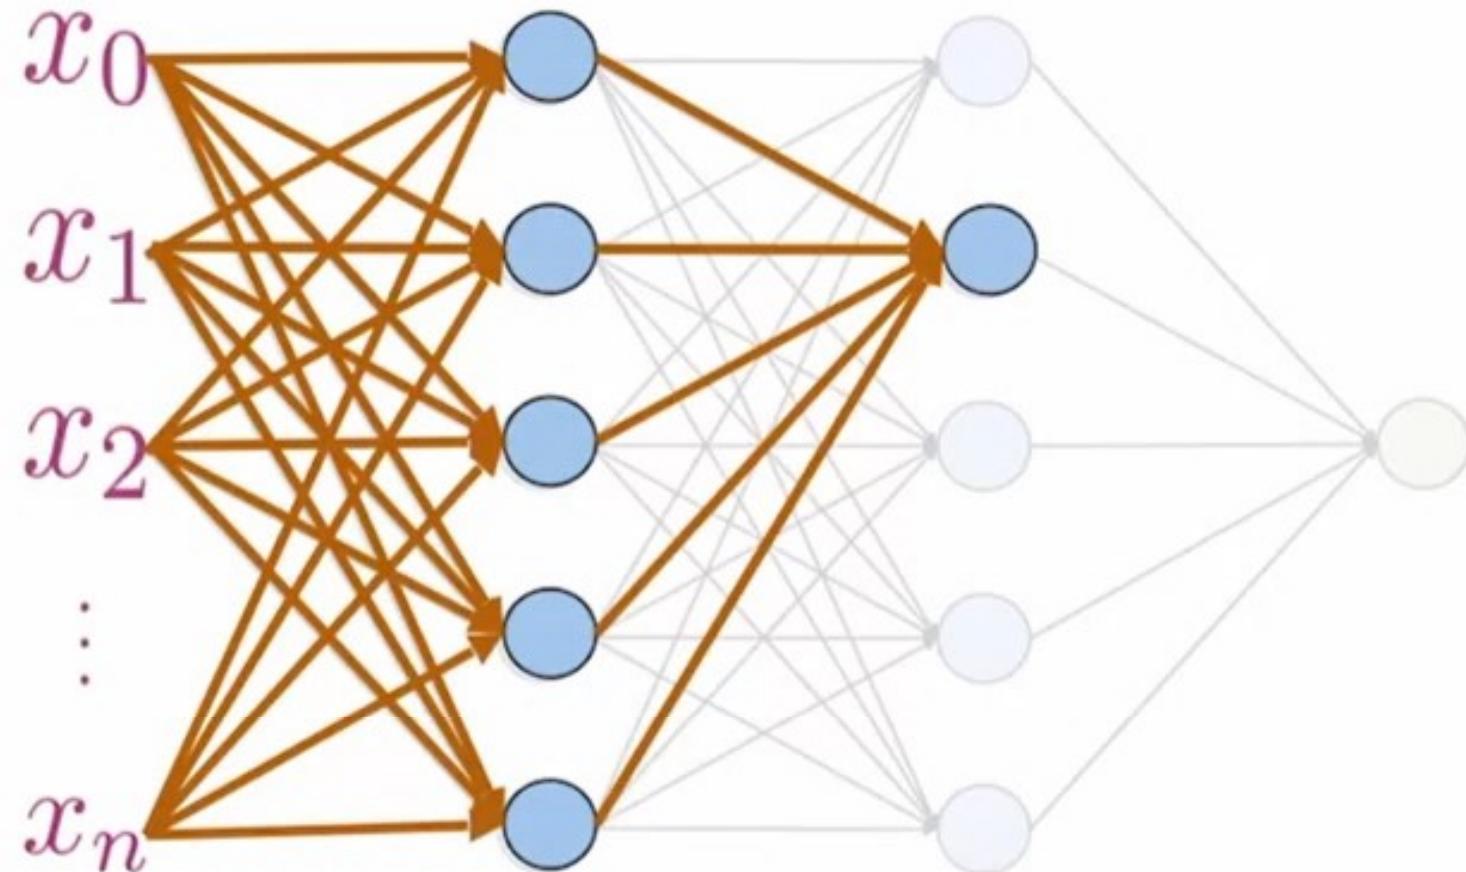


$$\hat{z}_i^{[l]} = \frac{z_i^{[l]} - \mu_{z_i^{[l]}}}{\sqrt{\sigma_{z_i^{[l]}}^2 + \epsilon}}$$

Shift factor  
Scale Factor

Learnable  
parameters to get  
the optimal dist.

## Batch Normalization: Test



During testing, mean and variance of training sets are used i.e. they are fixed for each test batches unlike training batches.

$$\hat{z}_i^{[l]} = \frac{z_i^{[l]} - \mathbf{E}(z_i^{[l]})}{\sqrt{\text{Var}(z_i^{[l]}) + \epsilon}}$$

Running mean and running std from training

Frameworks like  
Tensorflow and Pytorch  
keep track of them

# Summary

- Batch norm introduces learnable shift and scale factors
- During test, the running statistics from training are used
- Frameworks take care of the whole process



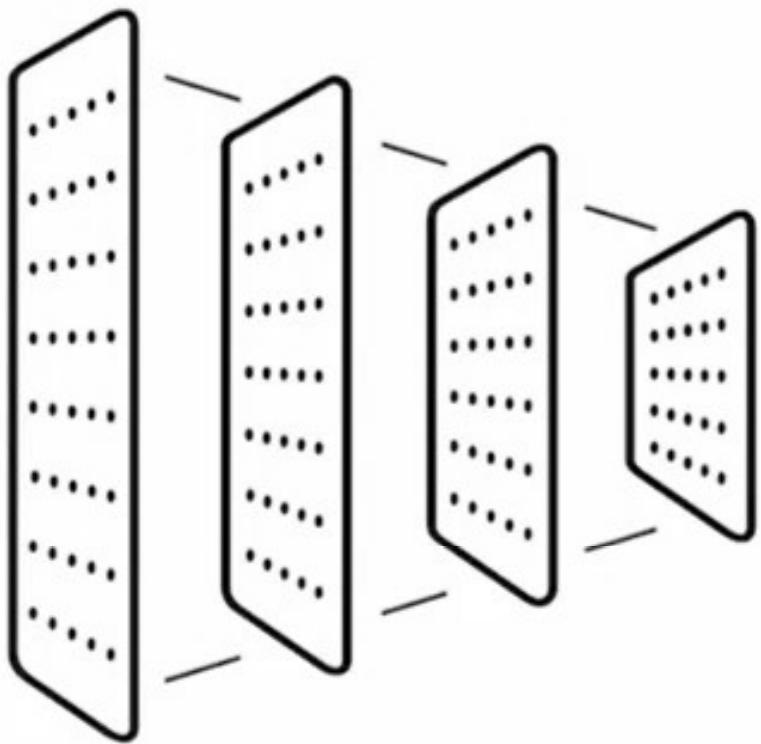


deeplearning.ai

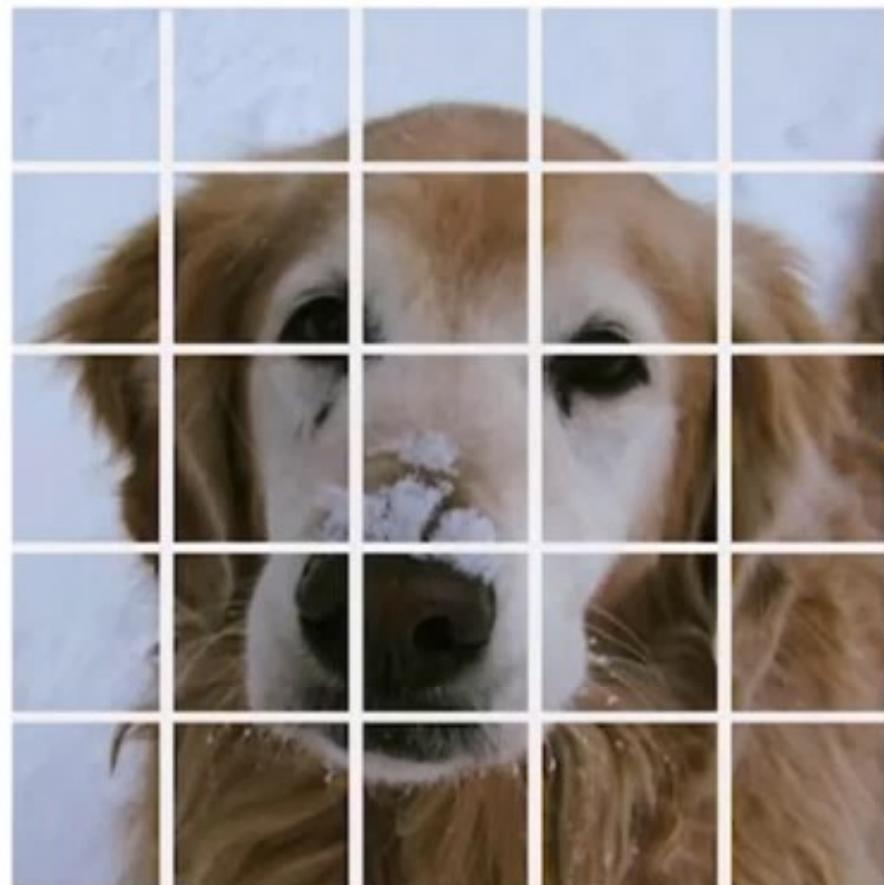
# Review of Convolutions

# Outline

- What convolutions are
- How they work



# What is a convolution?



Eye Filter

Nose filter

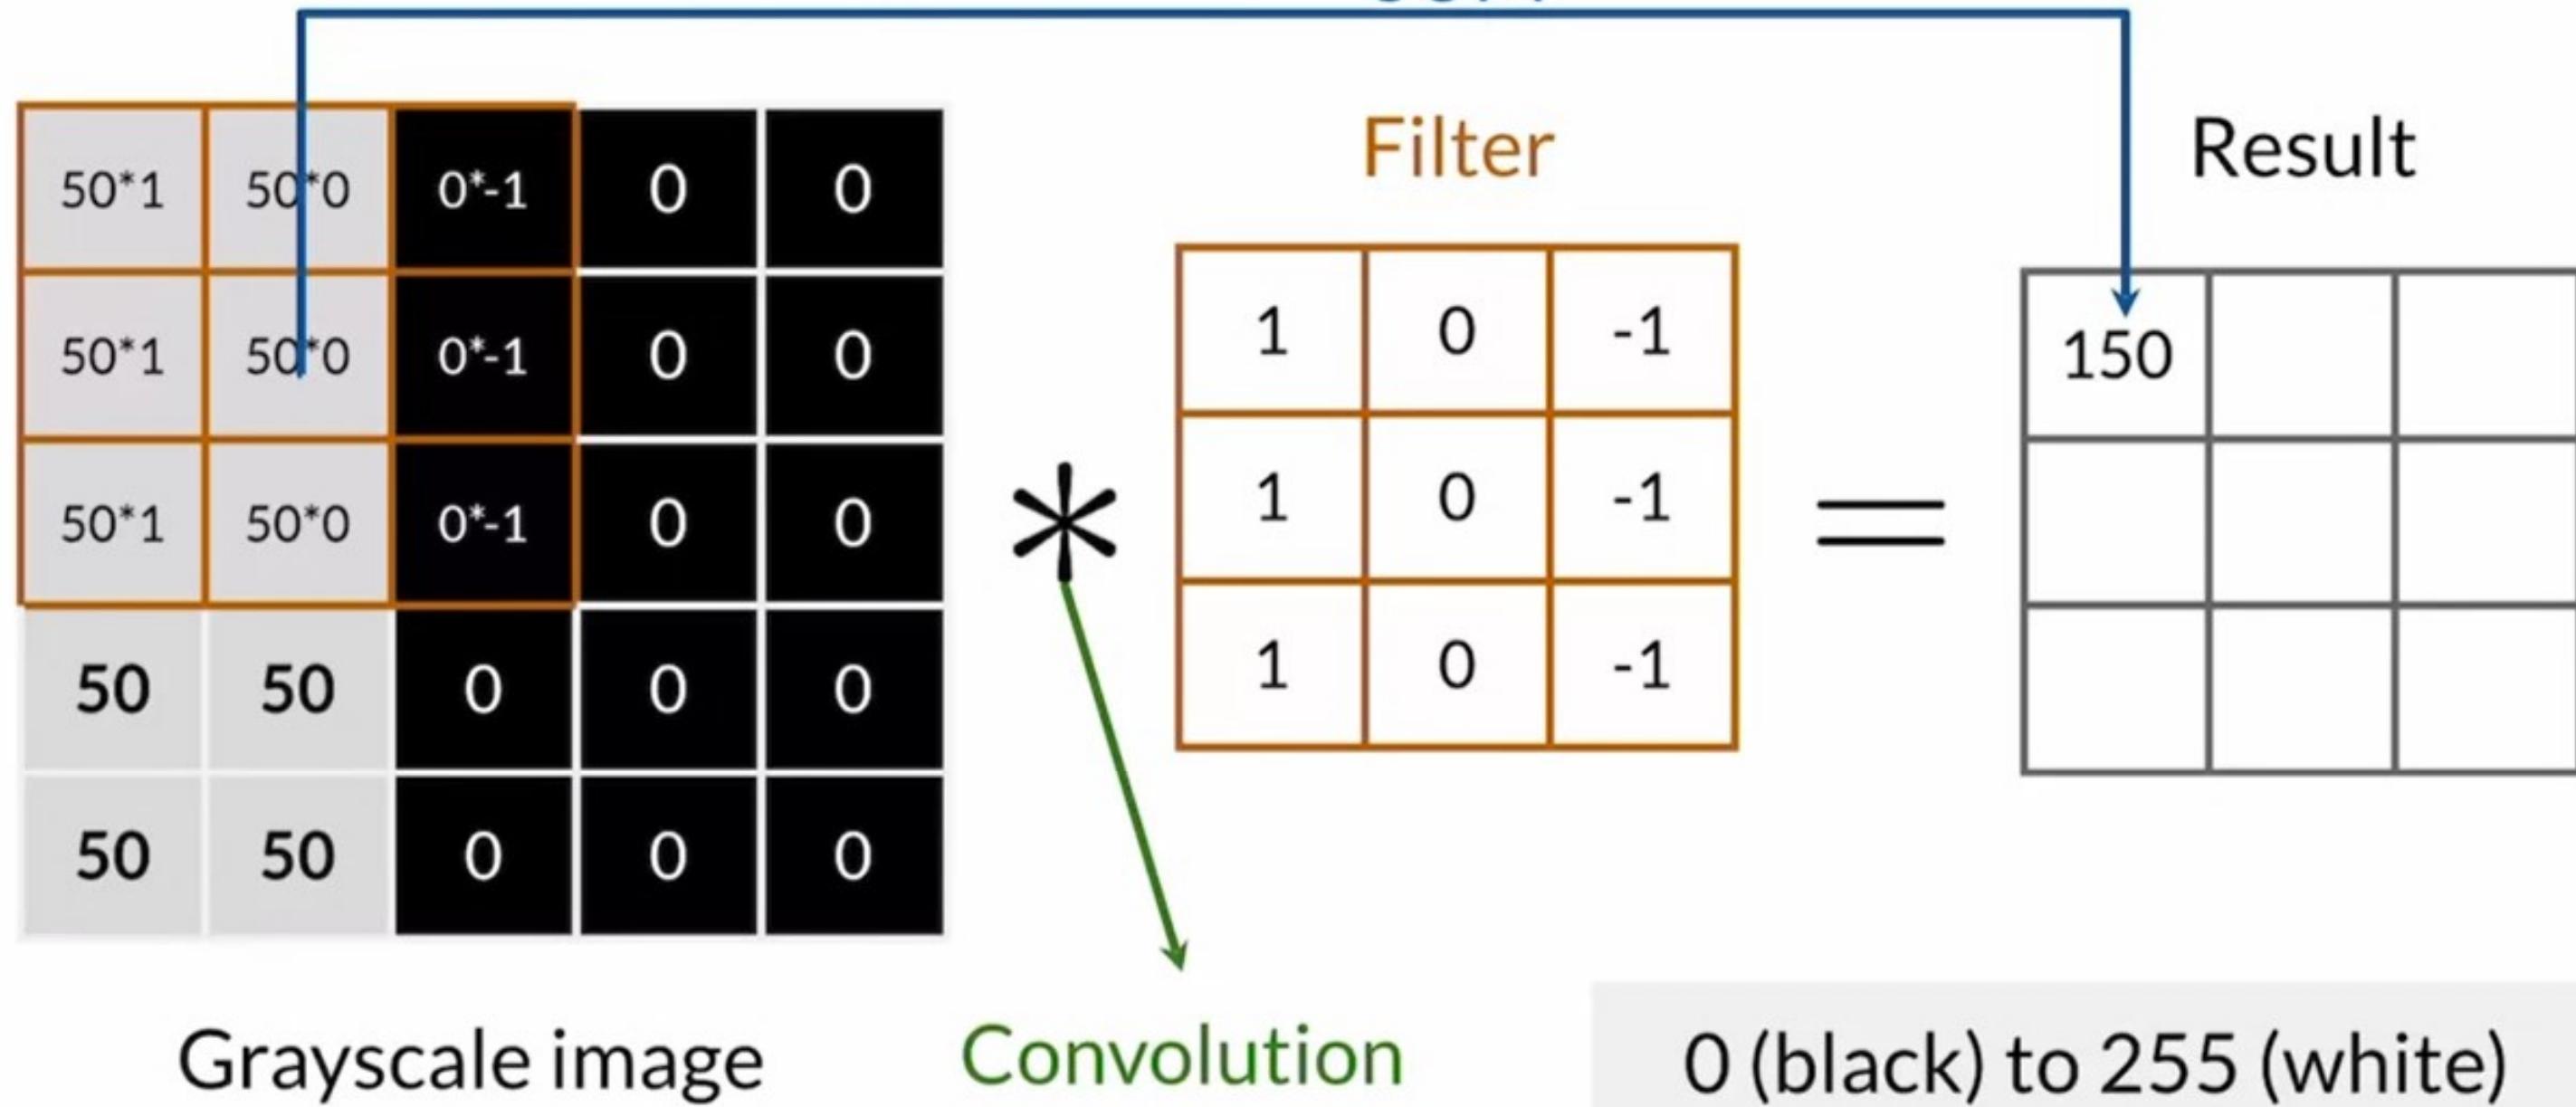
Ear filter

5	3	1
10	23	1
1	42	4

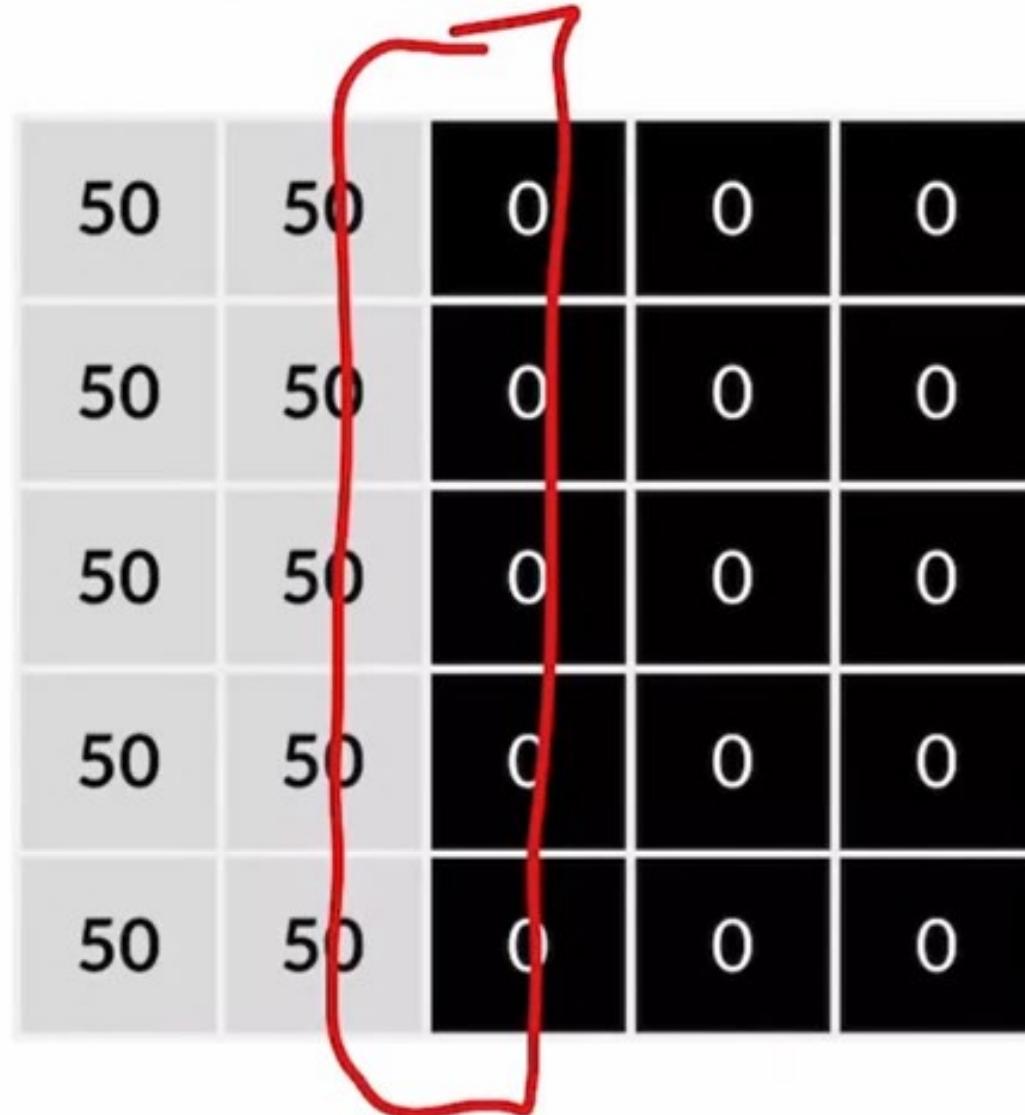
32	2	24
25	12	66
3	45	2

2	21	5
23	45	12
3	32	4

# What is a convolution?



# What is a convolution?



Grayscale image



Convolution

1	0	-1
1	0	-1
1	0	-1

Filter

Vertical edge detector

Result

150	150	0
150	150	0
150	150	0

=

0 (black) to 255 (white)

# Summary

- Convolutions are useful layers for processing images
- They scan the image to detect useful features
- Just element-wise products and sums!



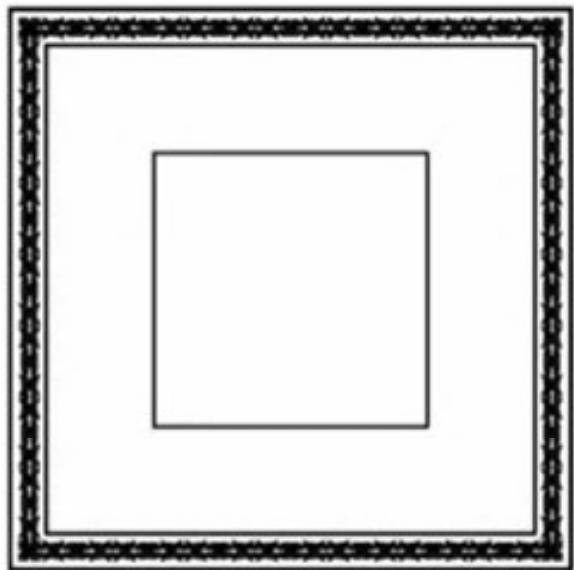


deeplearning.ai

# Padding and Stride

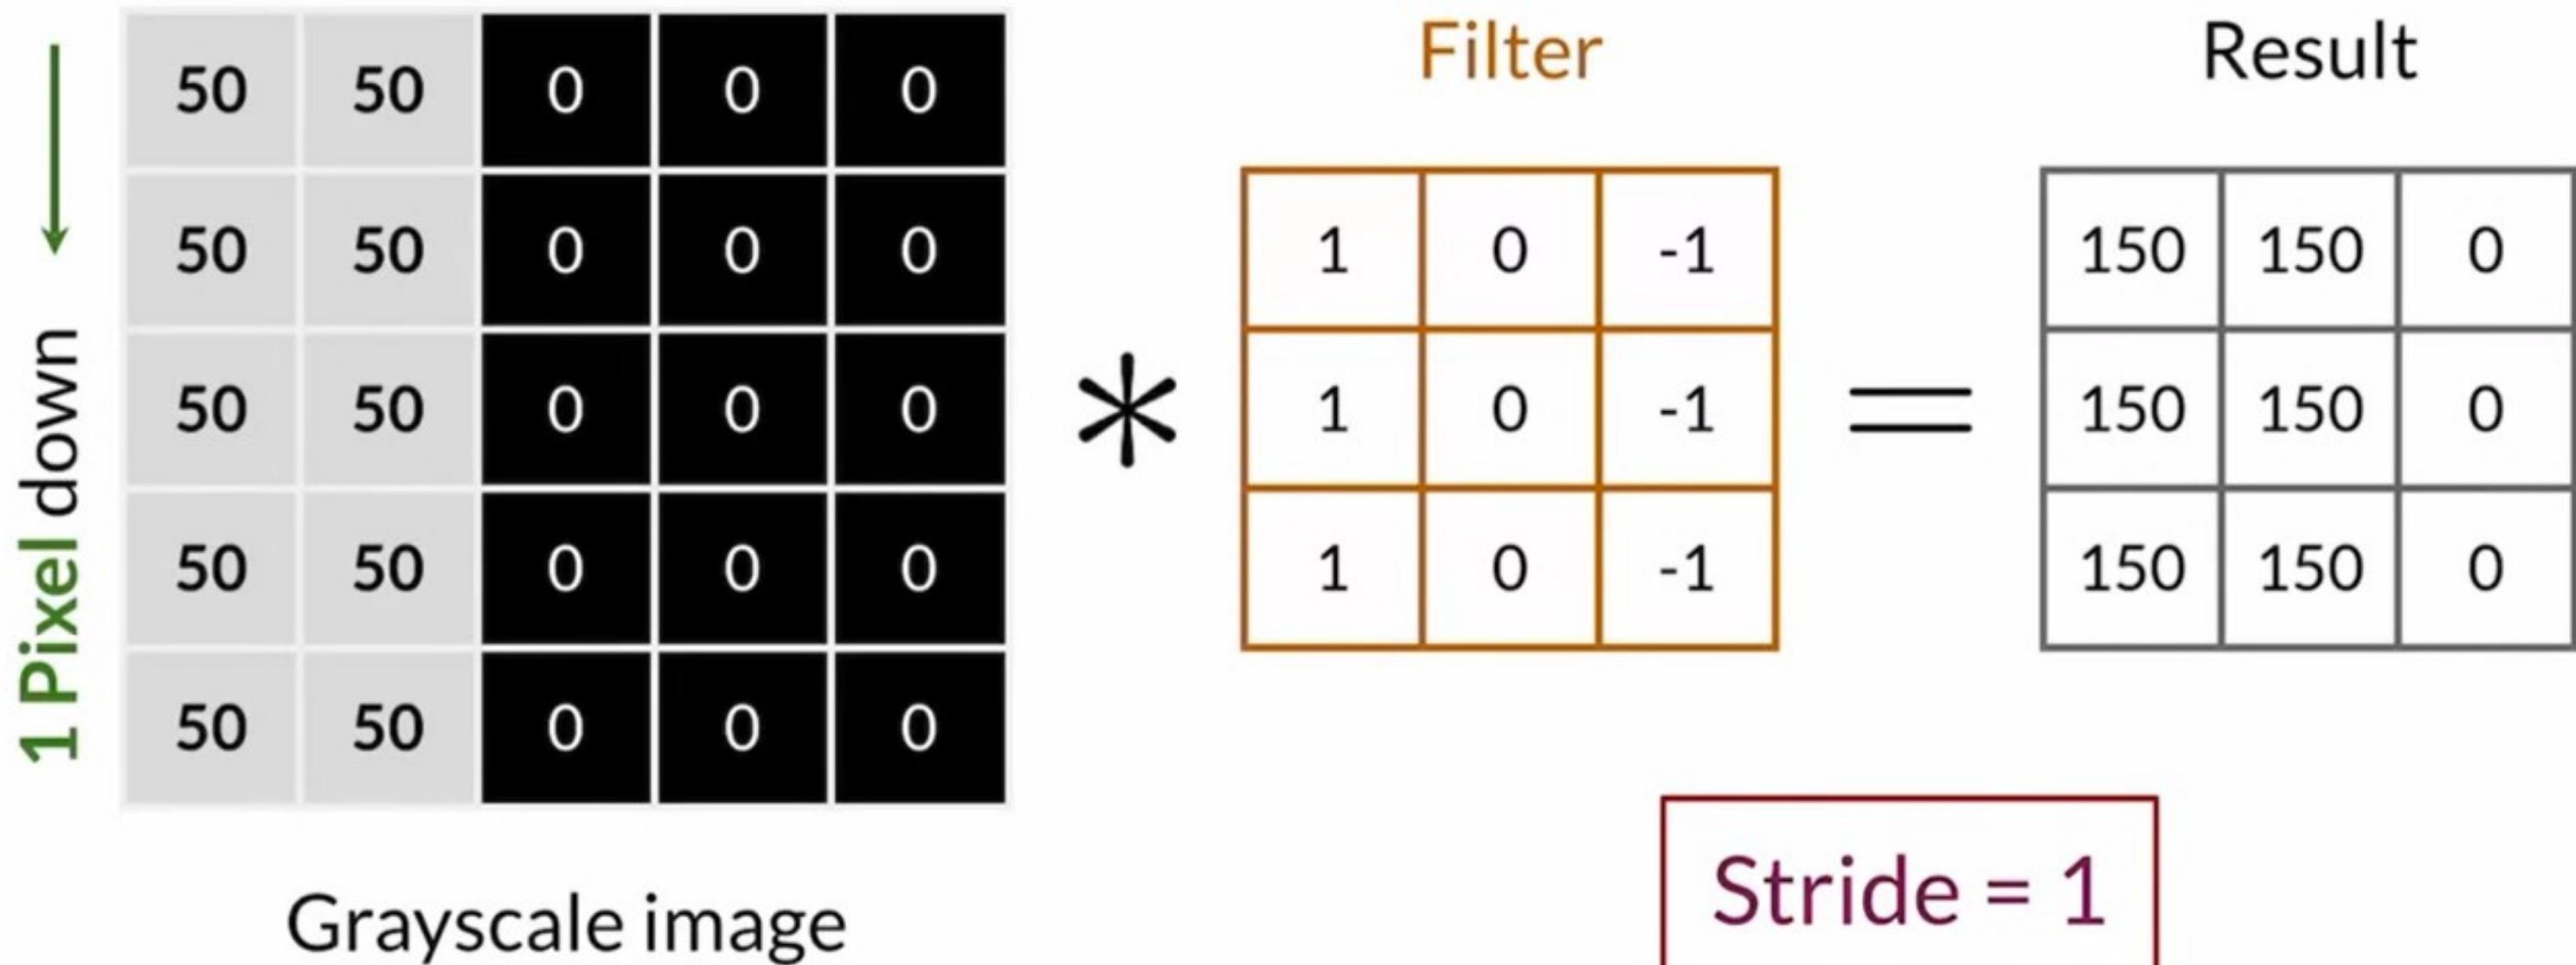
# Outline

- Padding and stride
- The intuition behind padding



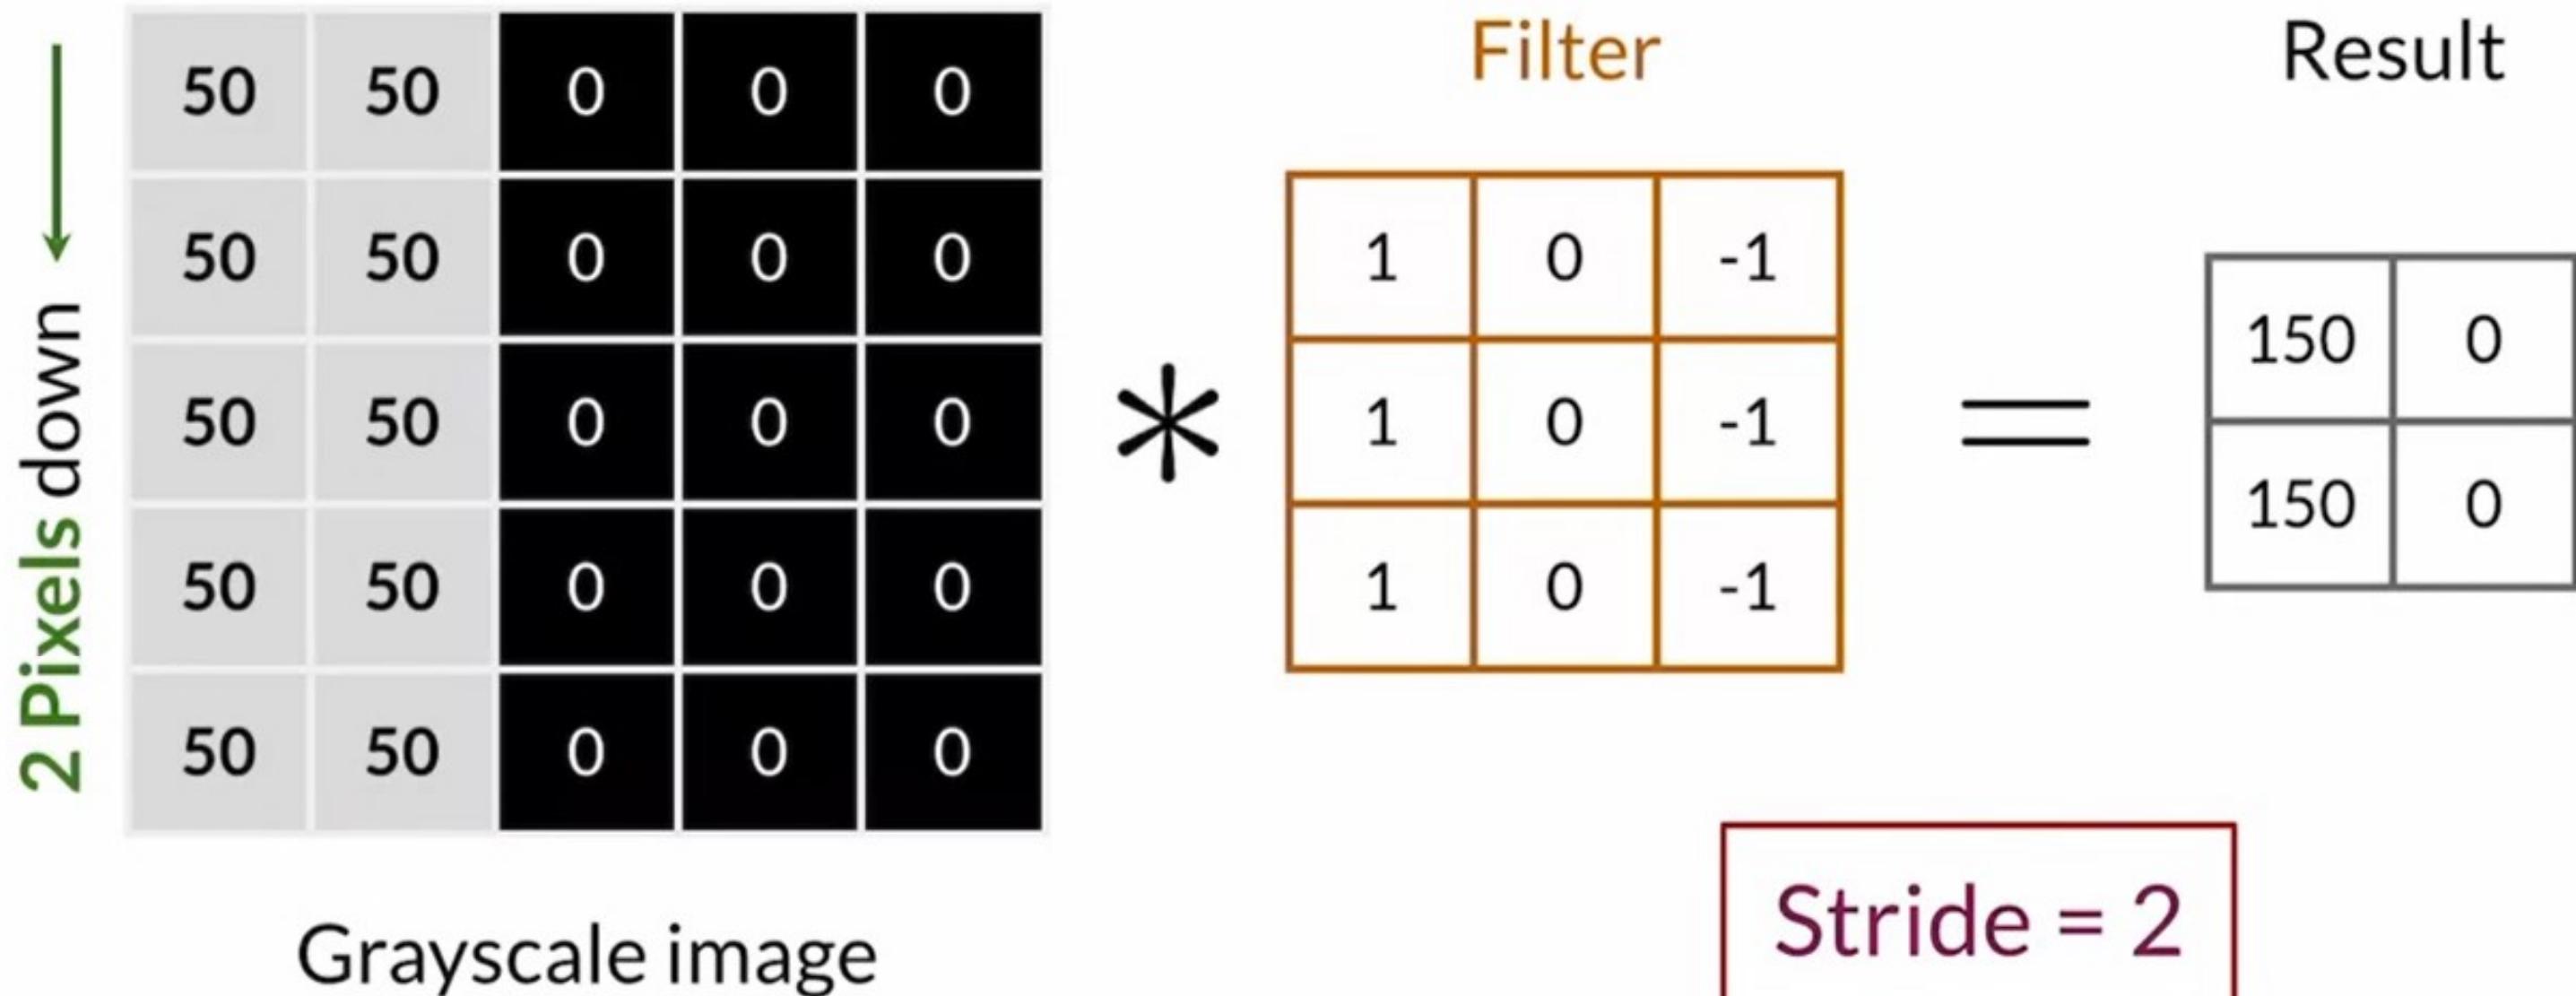
# Stride

→ 1 Pixel to the right

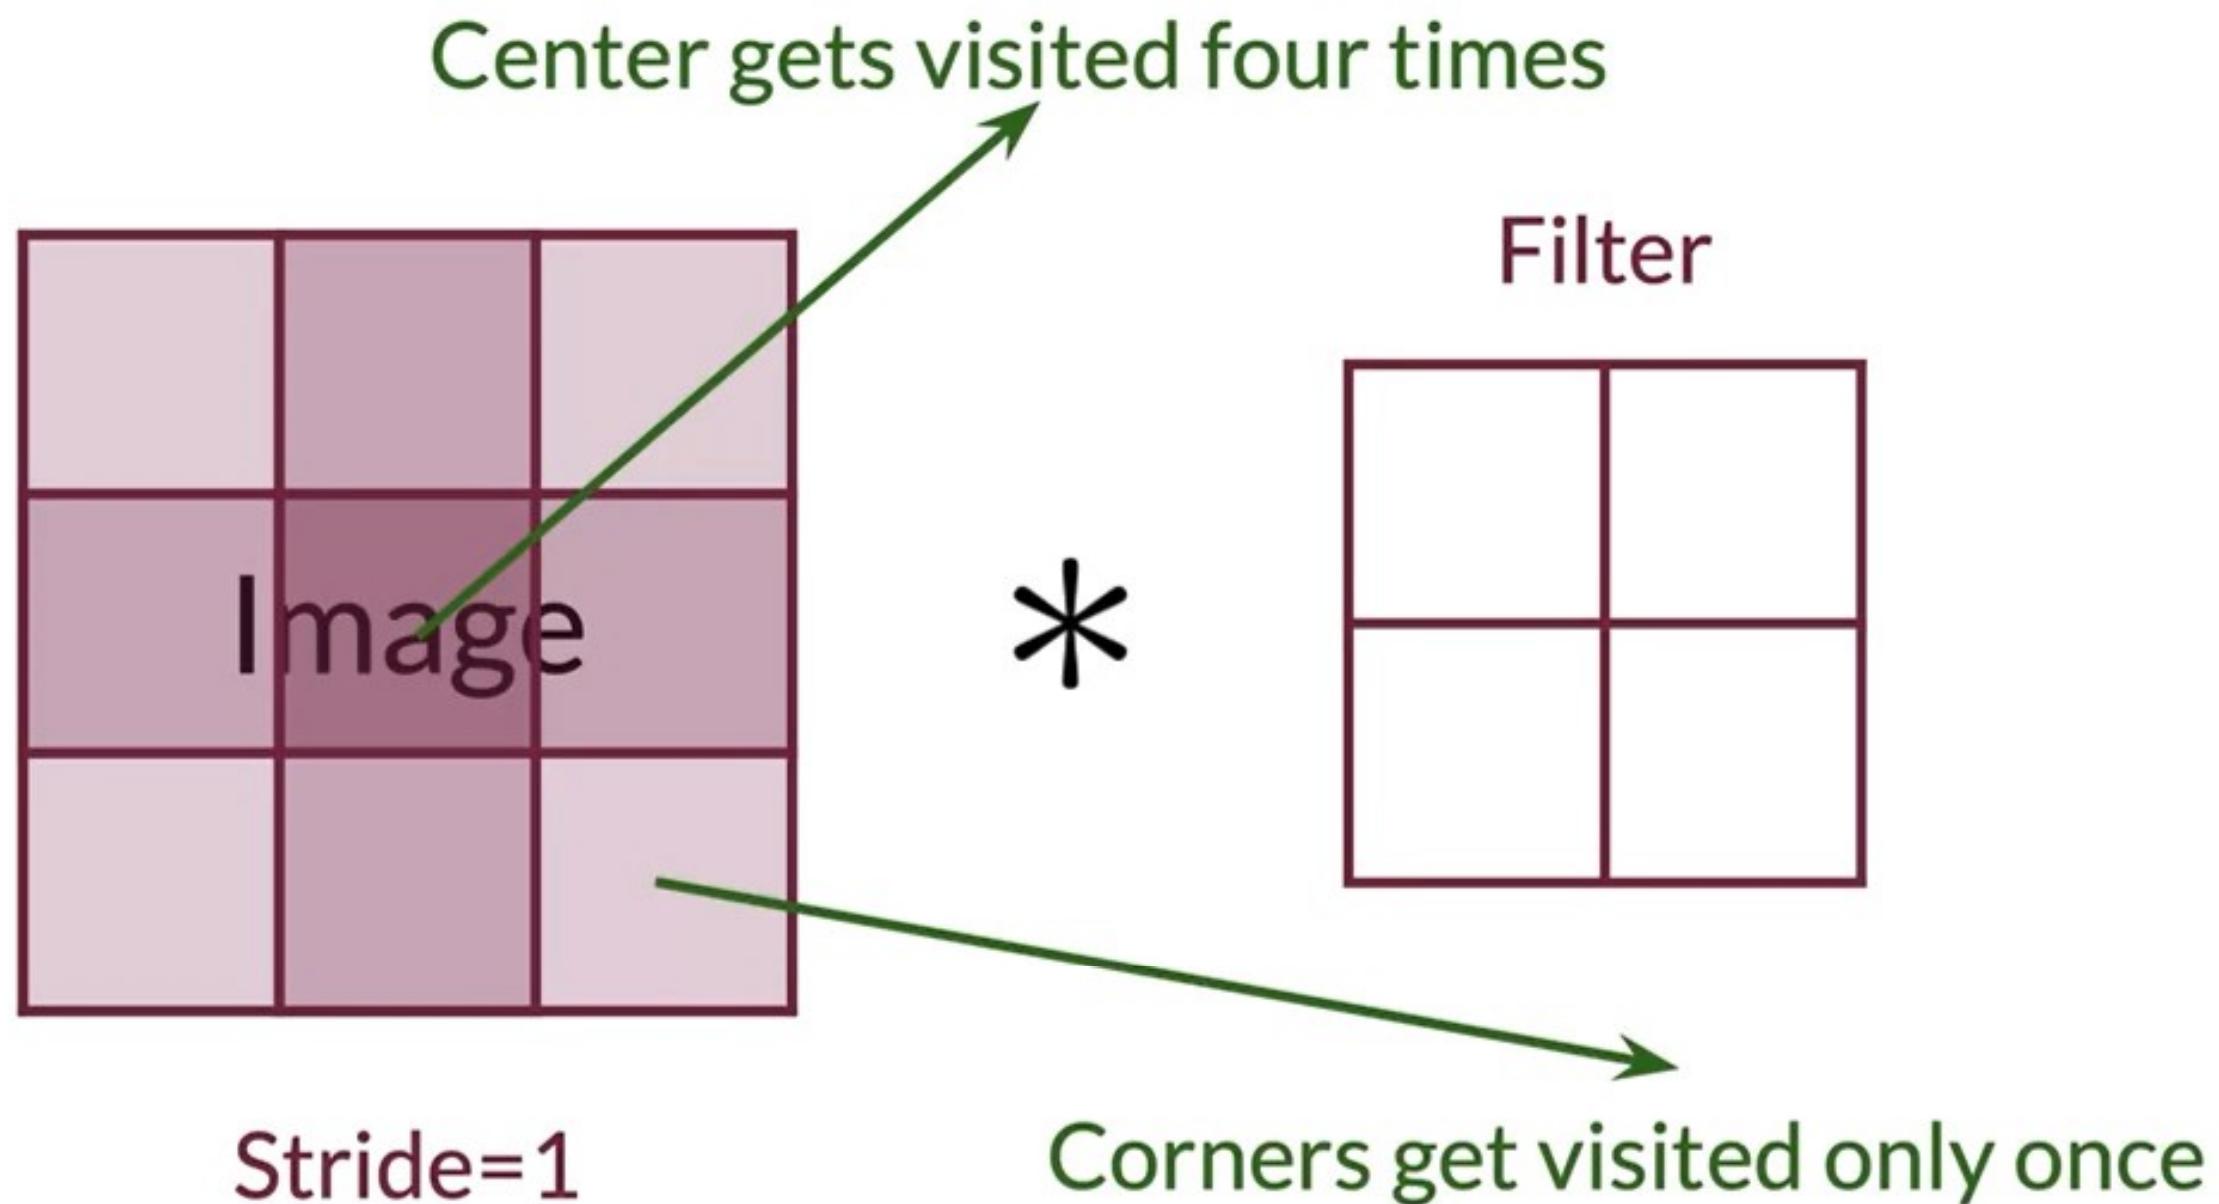


# Stride

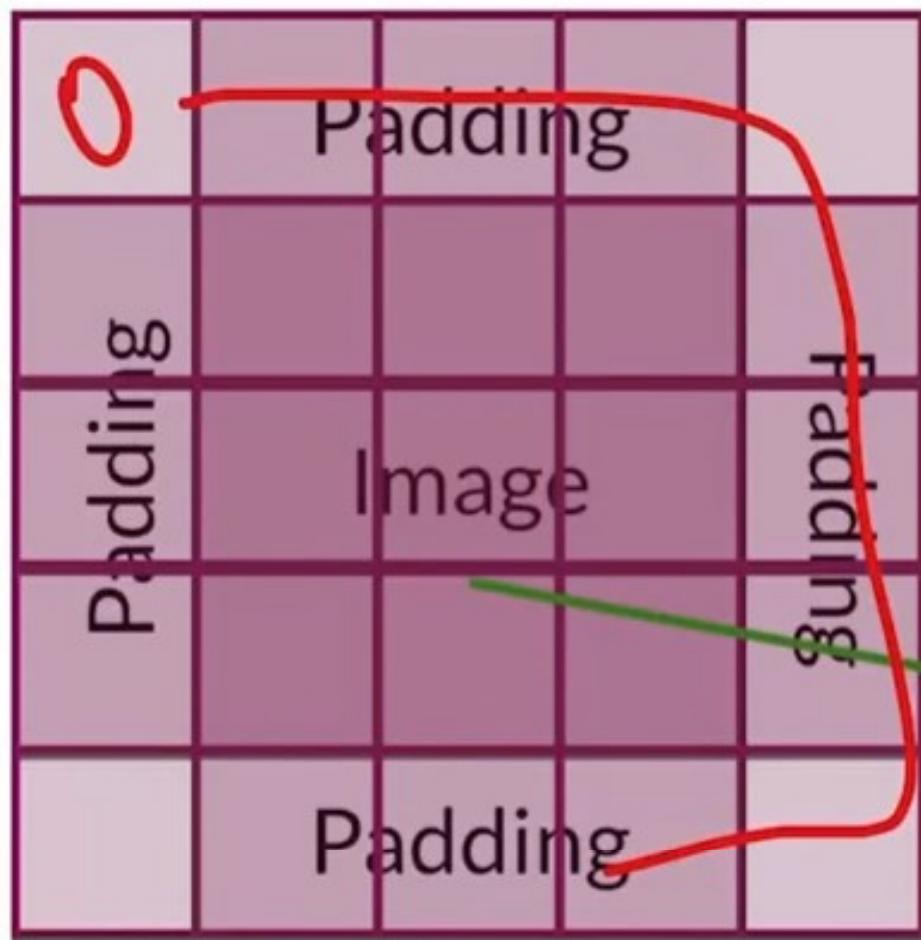
→ 2 Pixels to the right



# Padding



# Padding



Filter

\*



Every pixel within the image gets visited the same number of times

# Summary

- Stride determines how the filter scans the image
- Padding is like a frame on the image
- Padding gives similar importance to the edges and the center



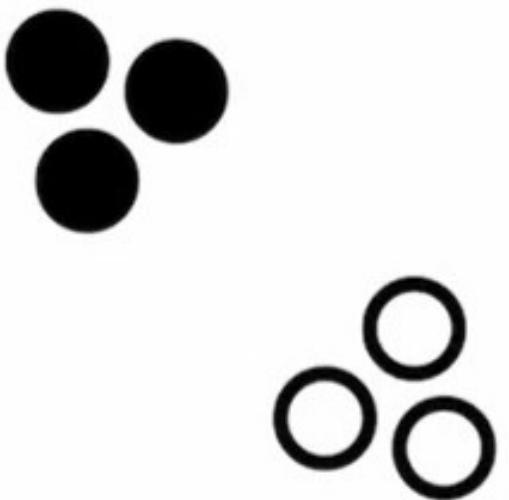


deeplearning.ai

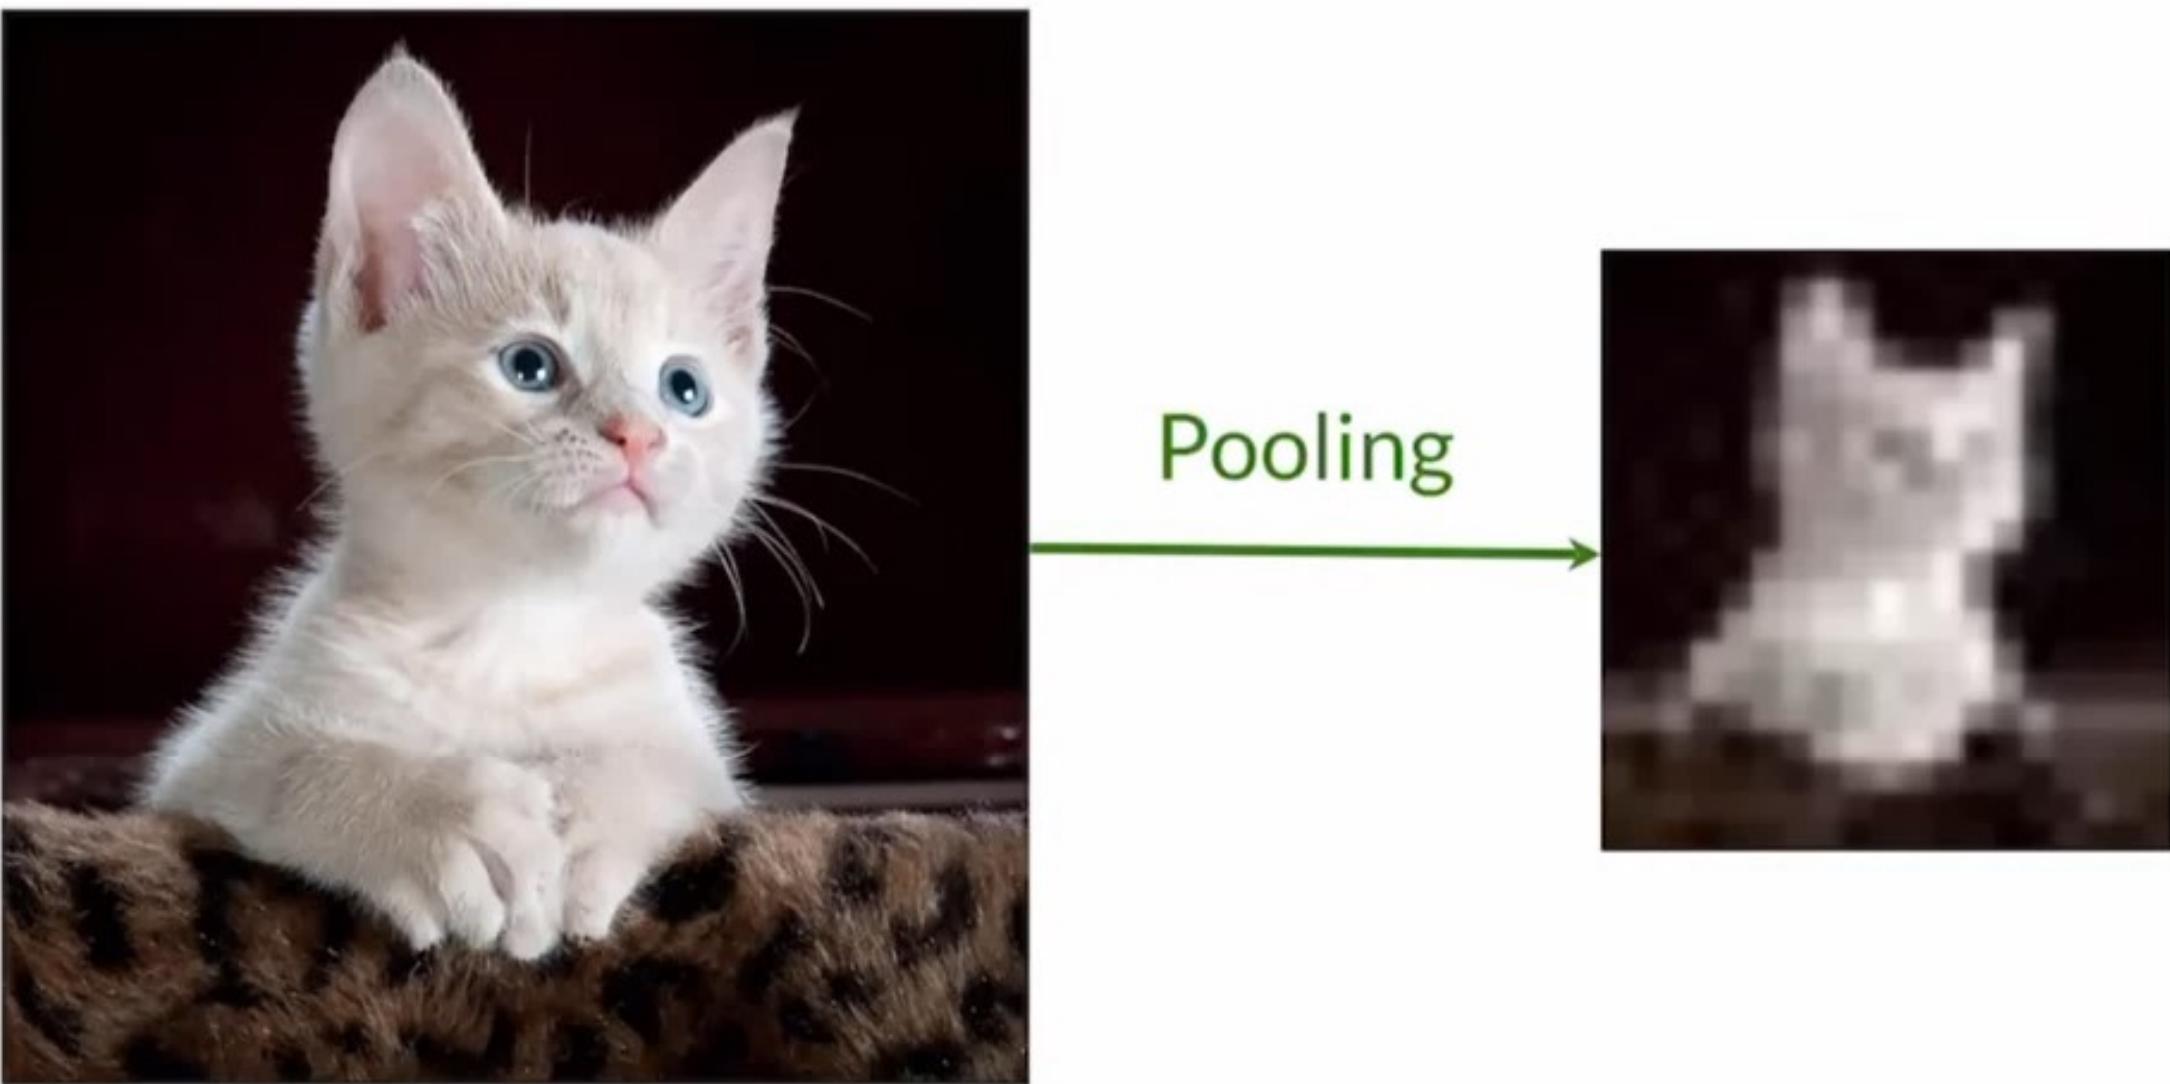
# Pooling and Upsampling

# Outline

- Pooling
- Upsampling and its relation to pooling



# Pooling



# Max Pooling

0	20	30	45
8	1	23	43
0	4	40	20
10	6	43	42

2x2 pooling with stride = 2

Max pooling

20	45
10	43

4x4 input to 2x2 output

# Max Pooling

NOTE: Pooling have no weights

0	20	30	45
8	1	23	43
0	4	40	20
10	6	43	42

4x4 input to 2x2 output

2x2 pooling with stride = 2

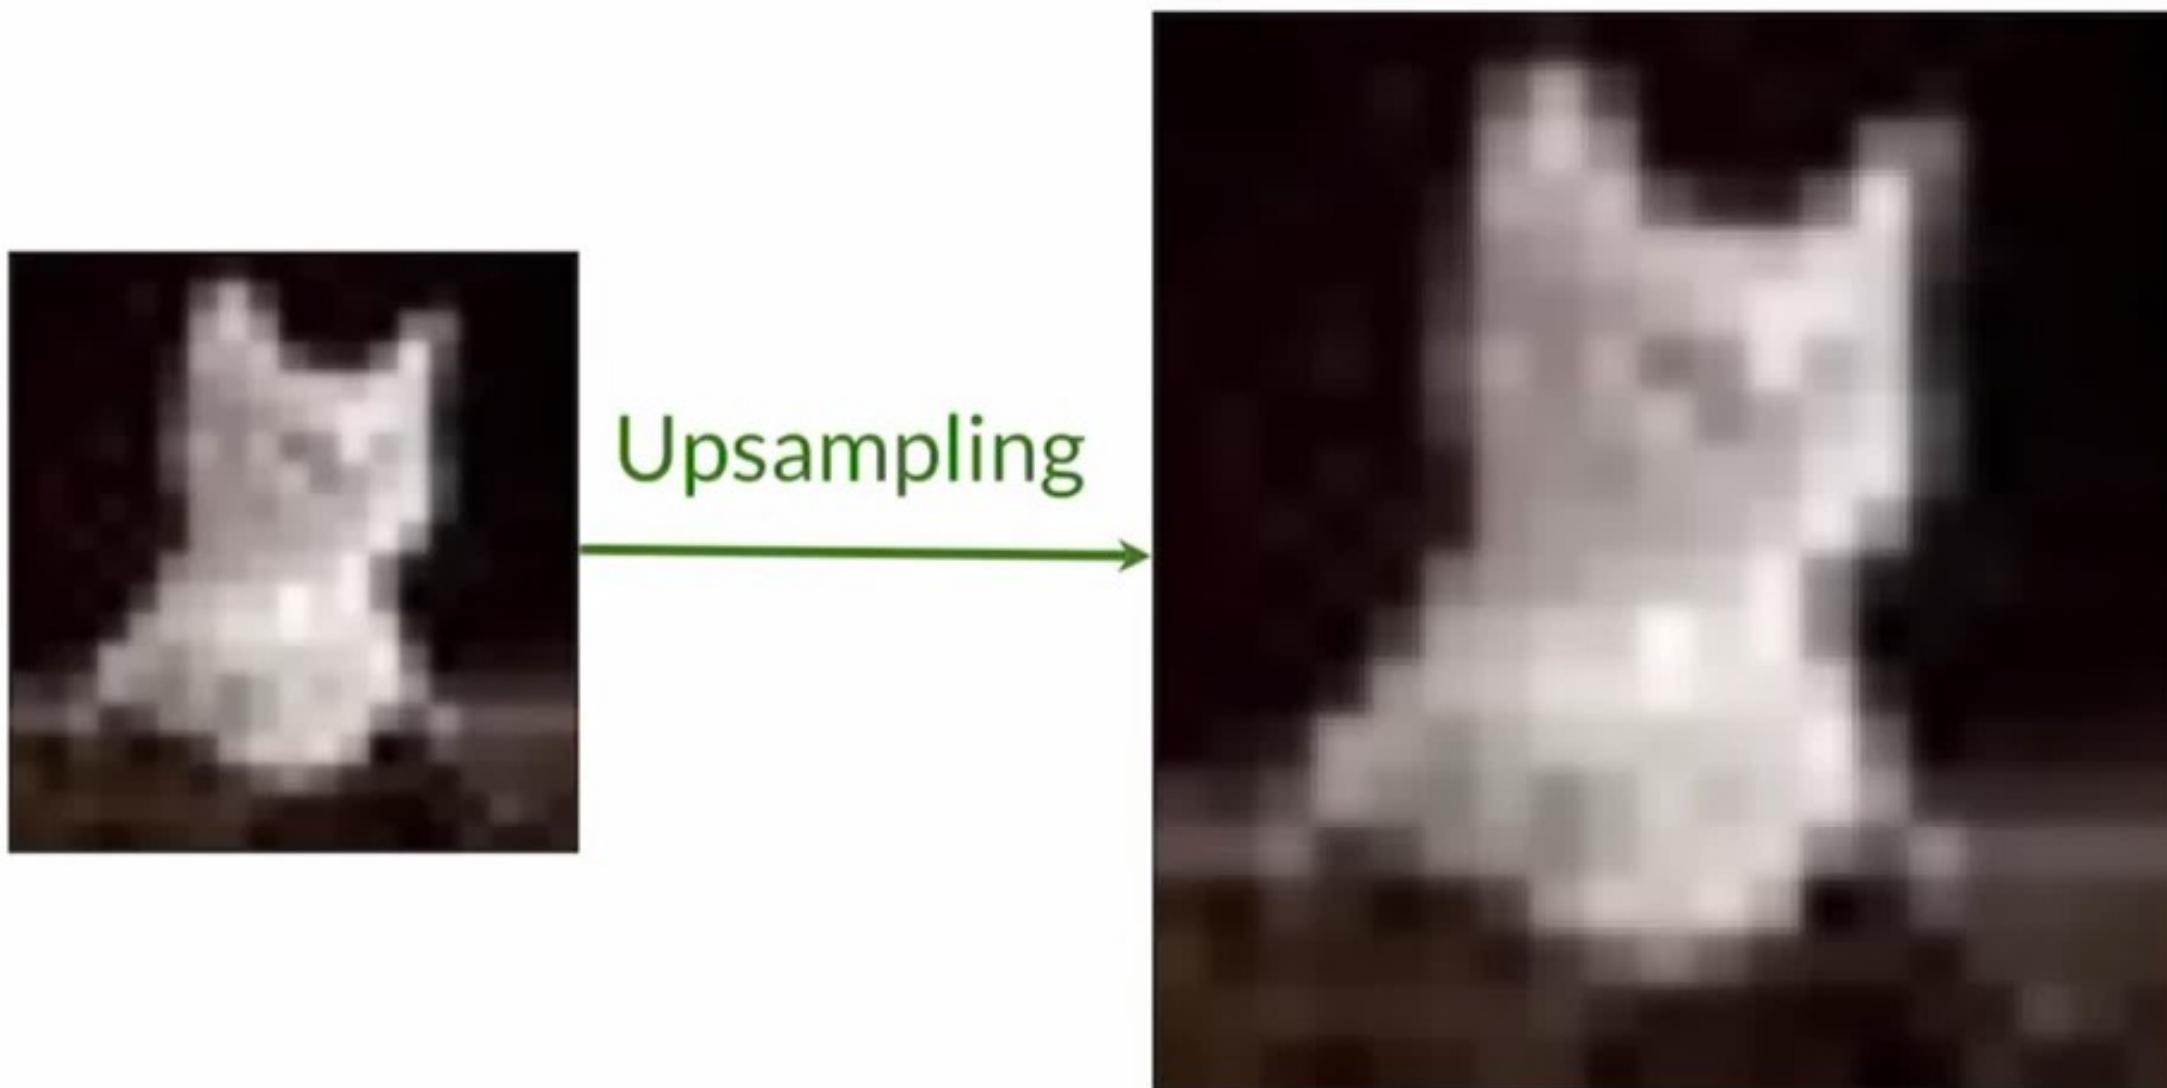
Max pooling

20	45
10	43

Other types include:

1. Average pooling
2. Min pooling

# Upsampling

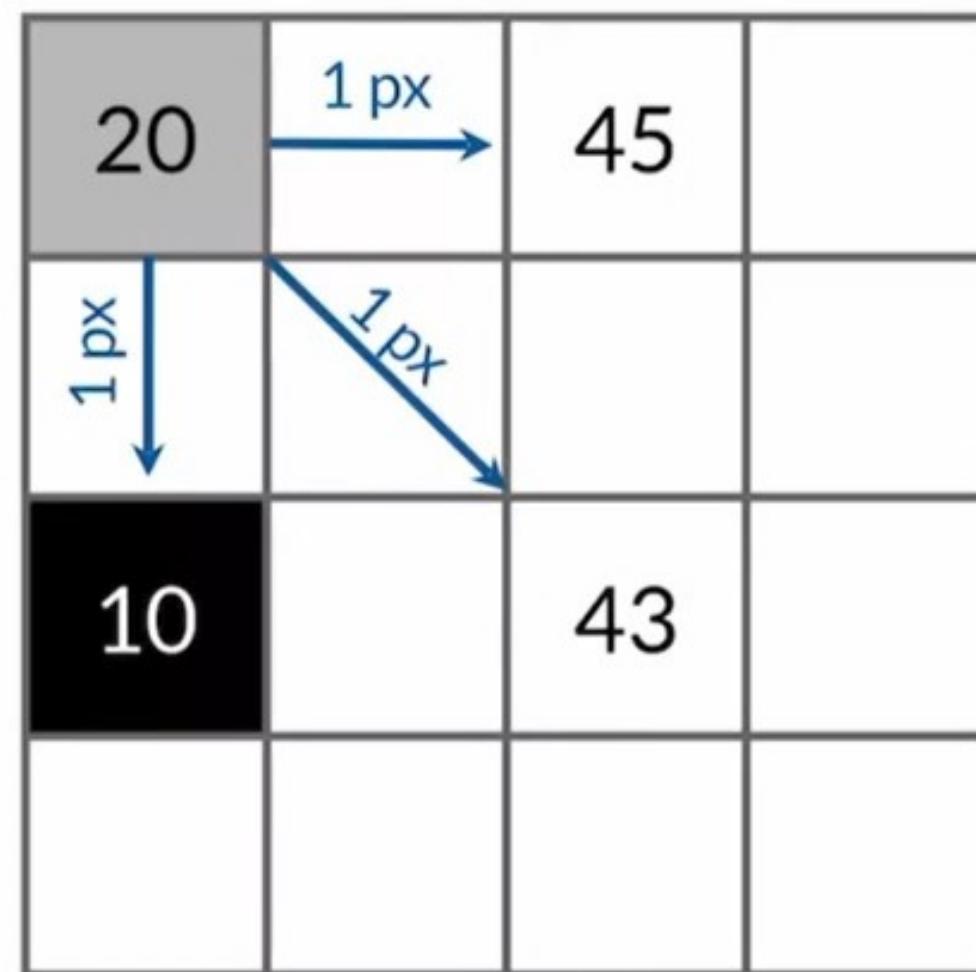


# Upsampling: Nearest Neighbors

2x2 input to 4x4 output

20	45
10	43

Upsampling



# Upsampling: Nearest Neighbors

2x2 input to 4x4 output

20	45
10	43

Upsampling

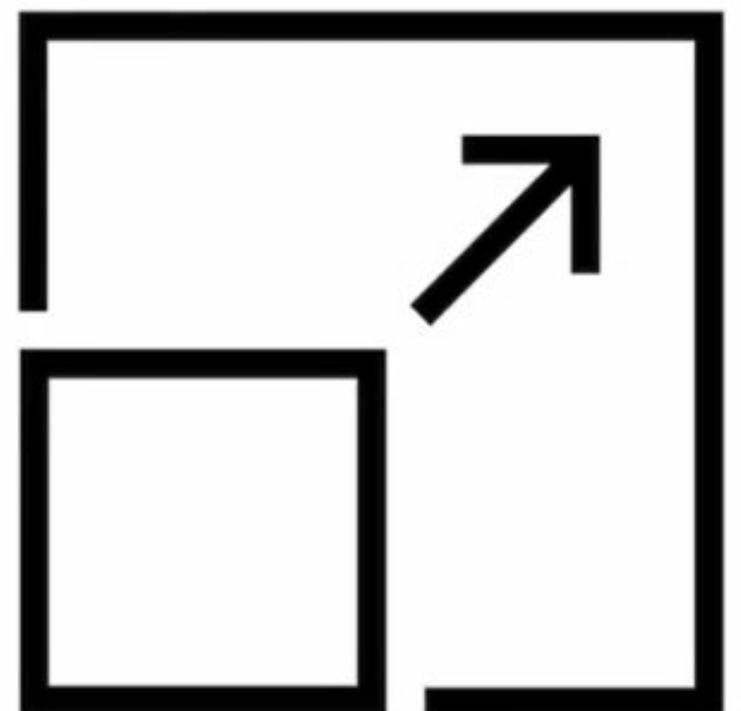
20	20	45	45
20	20	45	45
10	10	43	43
10	10	43	43

Other types include:

1. Linear interpolation
2. Bi-linear interpolation

# Summary

- Pooling reduces the size of the input
- Upsampling increases the size of the input
- No learnable parameters!



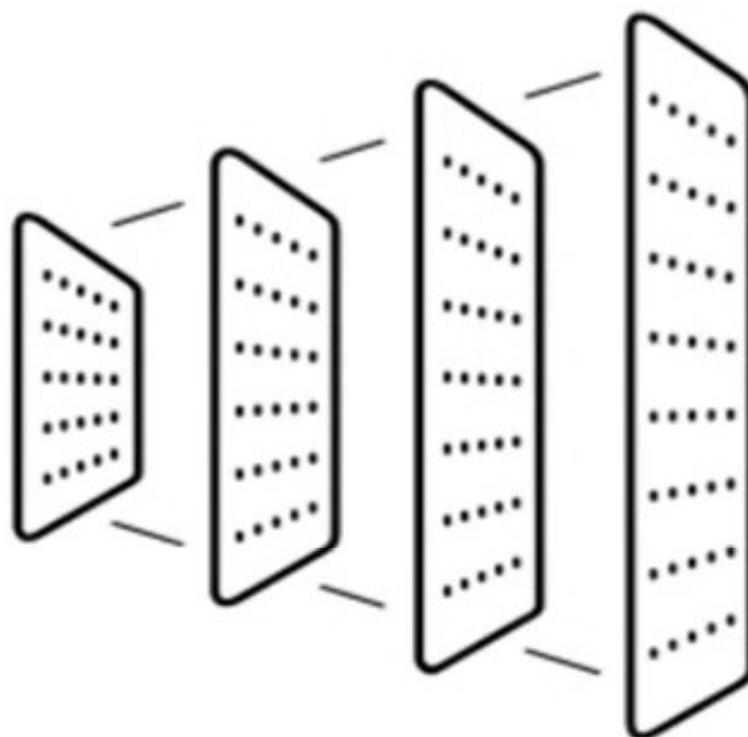


deeplearning.ai

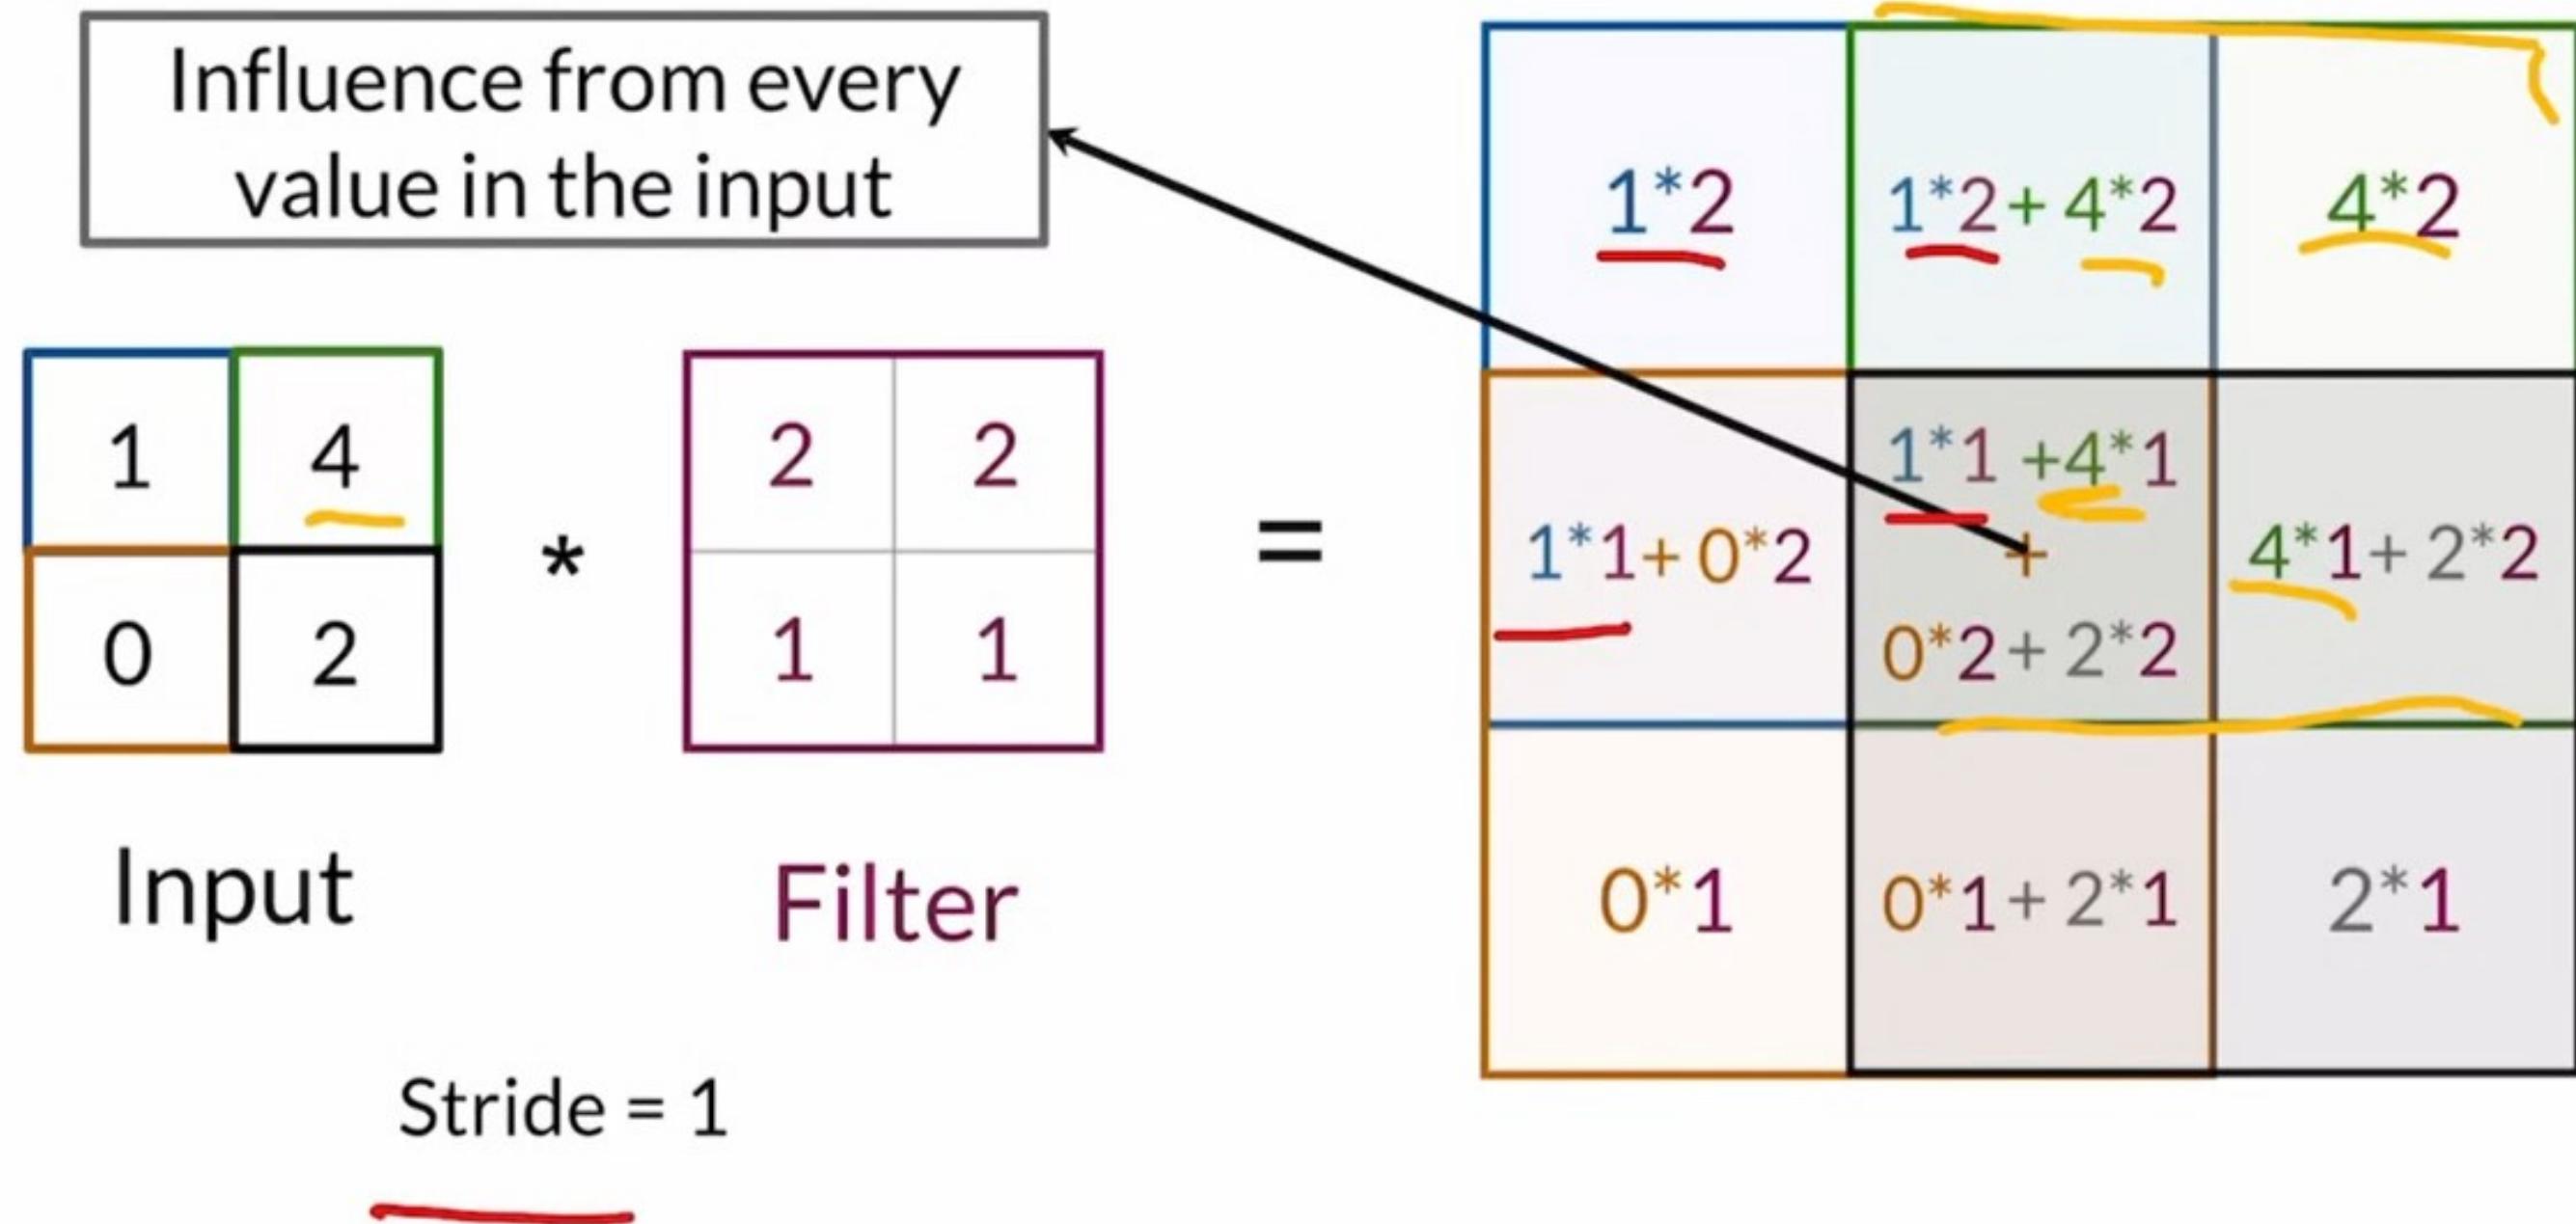
# Transposed Convolutions

# Outline

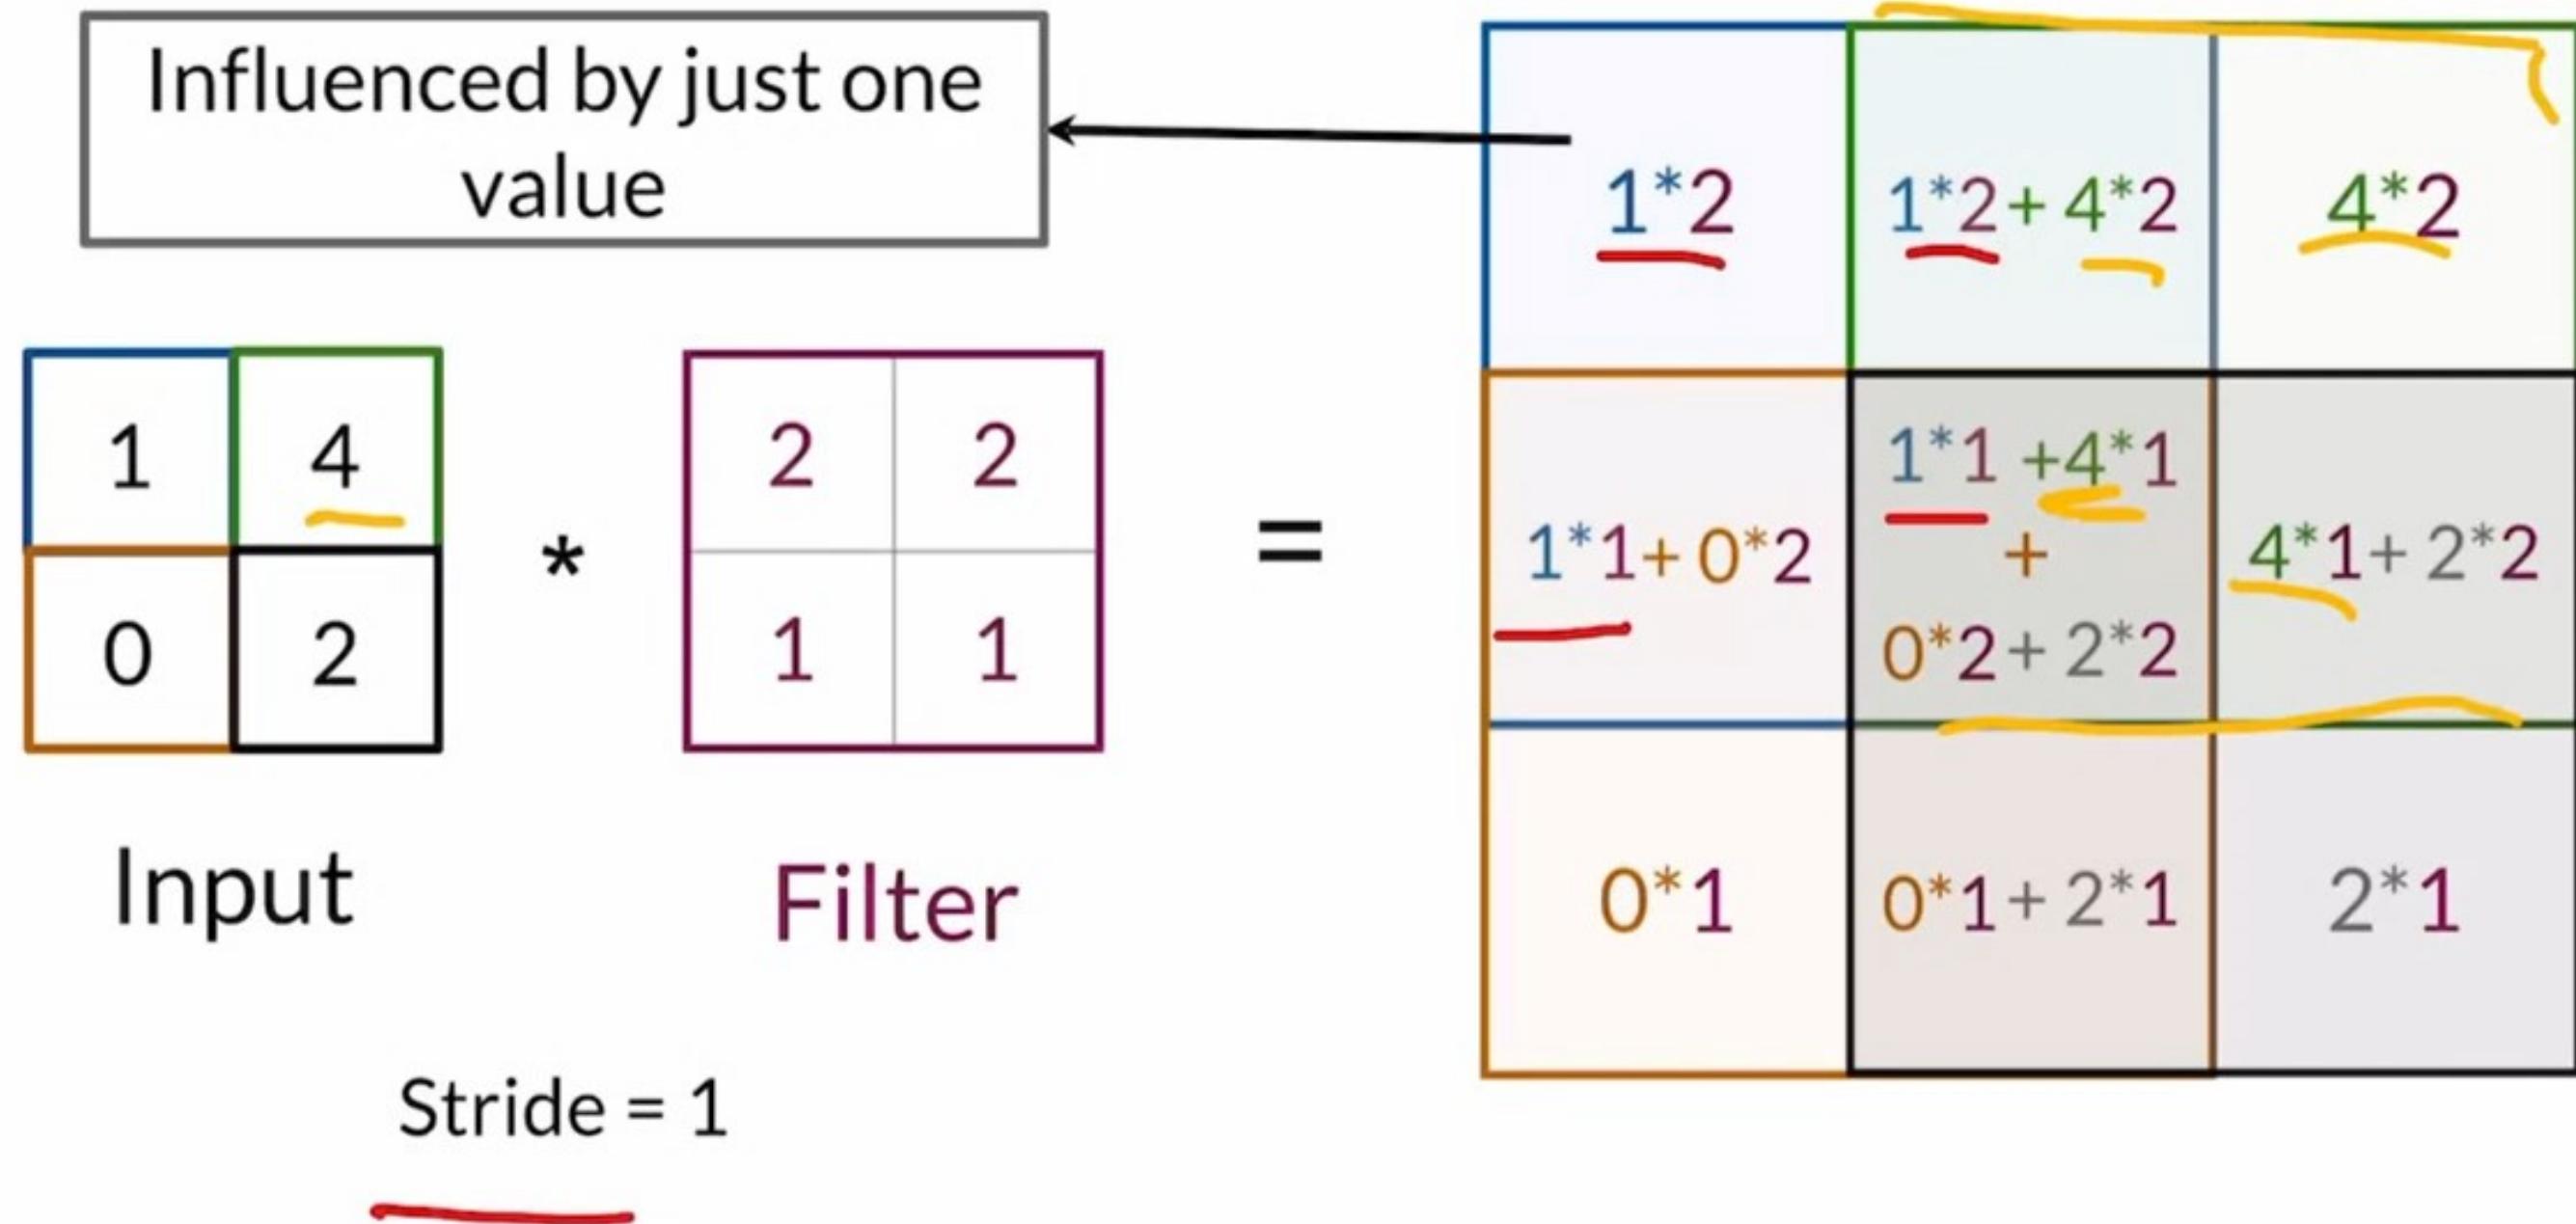
- Transposed convolutions as an upsampling technique
- Issues with transposed convolutions



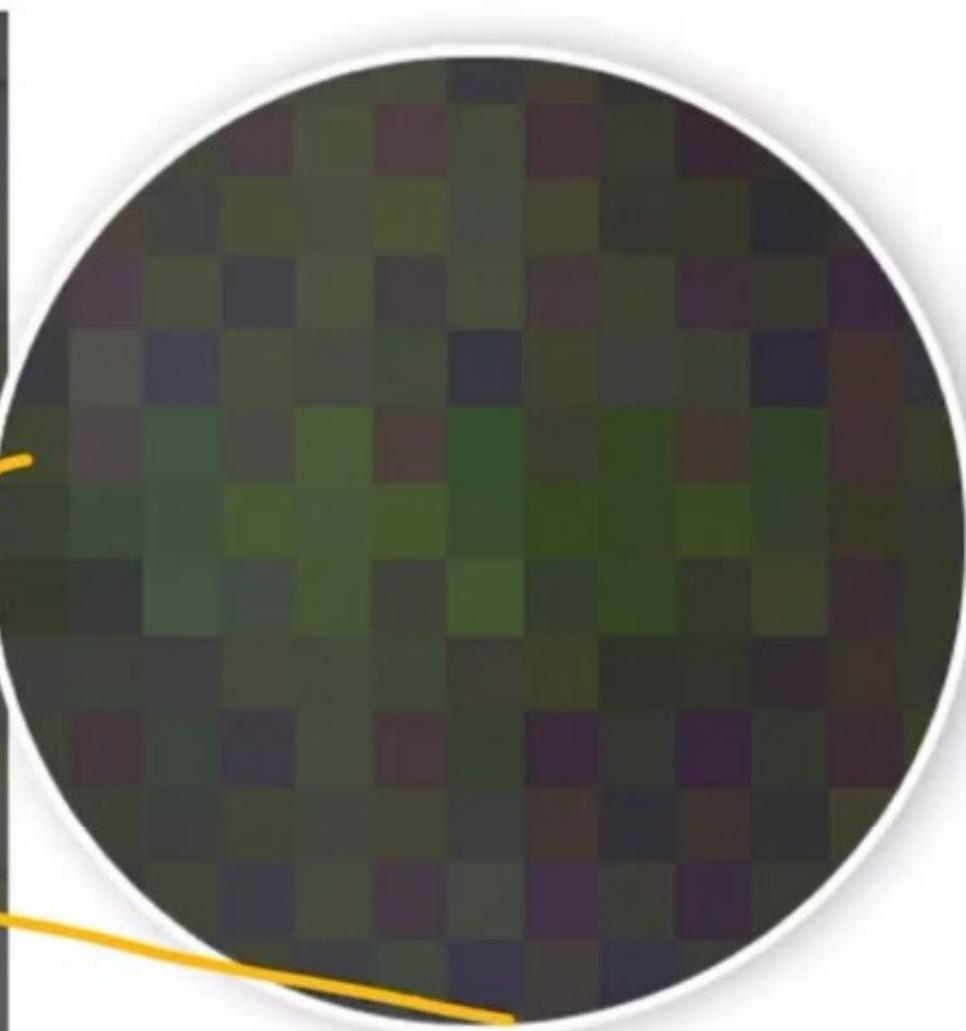
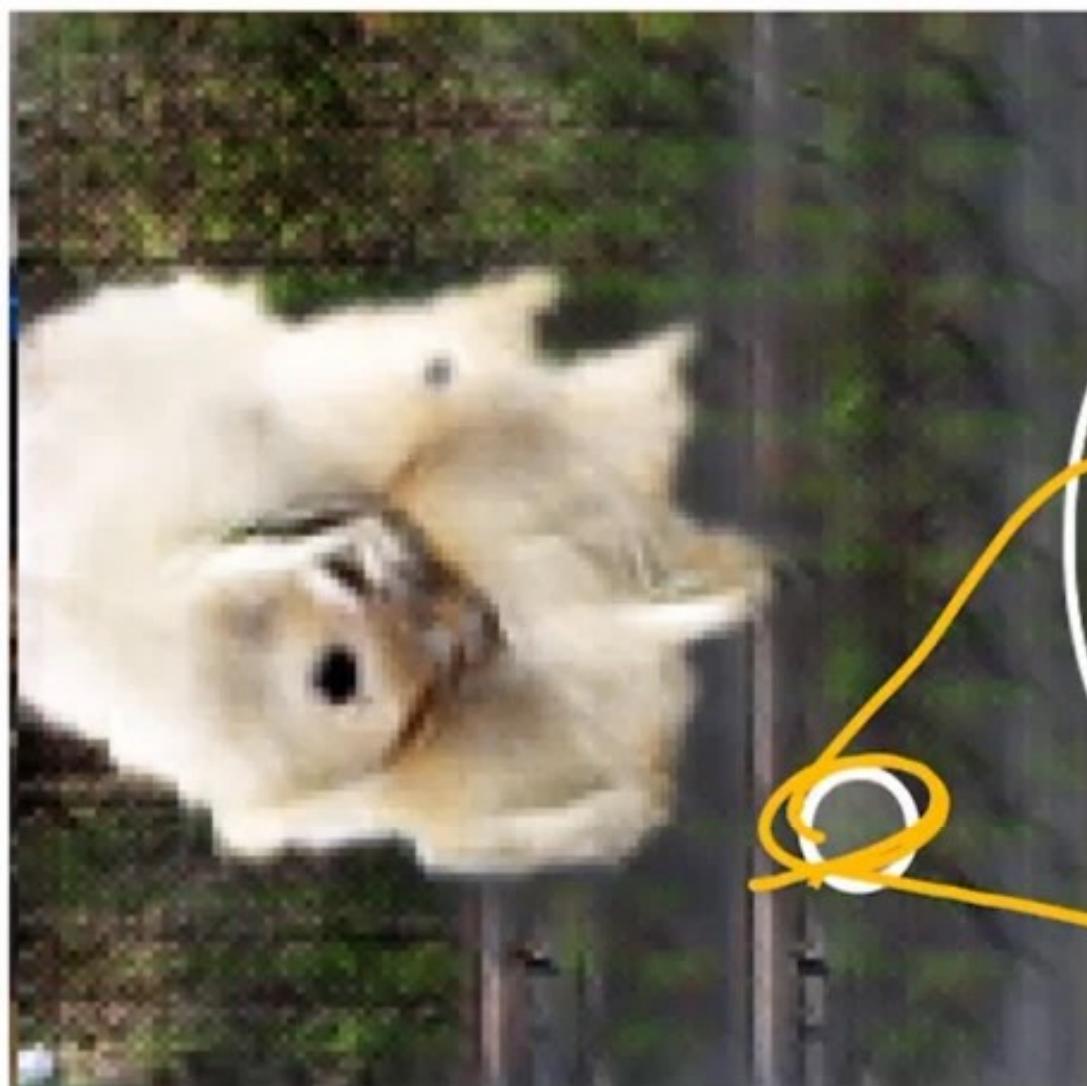
# Transposed Convolution



# Transposed Convolution



# The Problems with Transposed Convolution



Checkerboard  
Pattern

upsampling  
+  
convolution

Available from: <http://doi.org/10.23915/distill.00003>

# Summary

- Transposed convolutions upsample
- They have learnable parameters
- Problem: results have a checkerboard pattern

